




AWS Purpose Built Databases

Overview

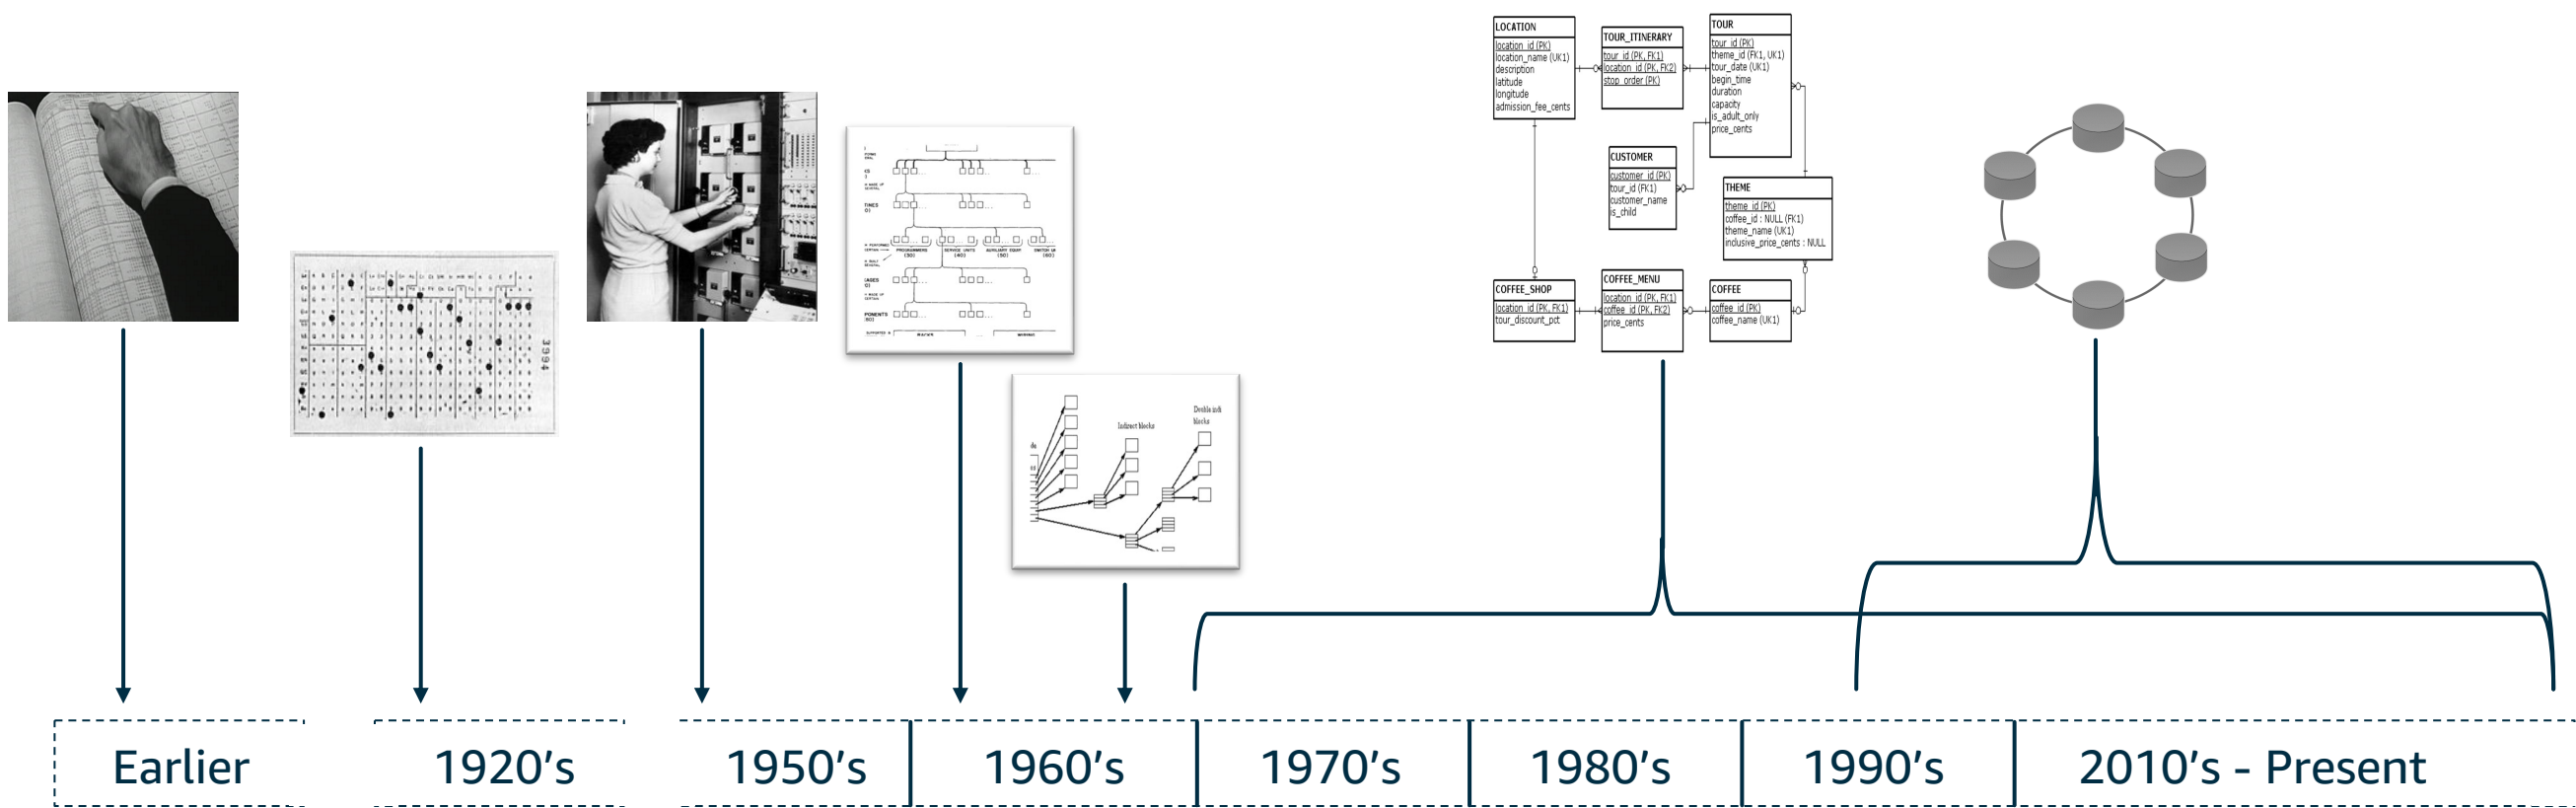
Eren Akbaba, AWS
eakbaba@amazon.com
Bilkent University




History of Databases



Timeline of Database Systems




Typical on-premise / monolithic architecture



Challenges faced with on-premise / monolithic architectures

Managing on-premise database environments - time consuming and complex




Scaling Limitations

- Scaling for peak, forecasting
- Cost of provisioning for peak and licensing
- Horizontally scaling /Sharding requires complex application logic



Availability Challenges

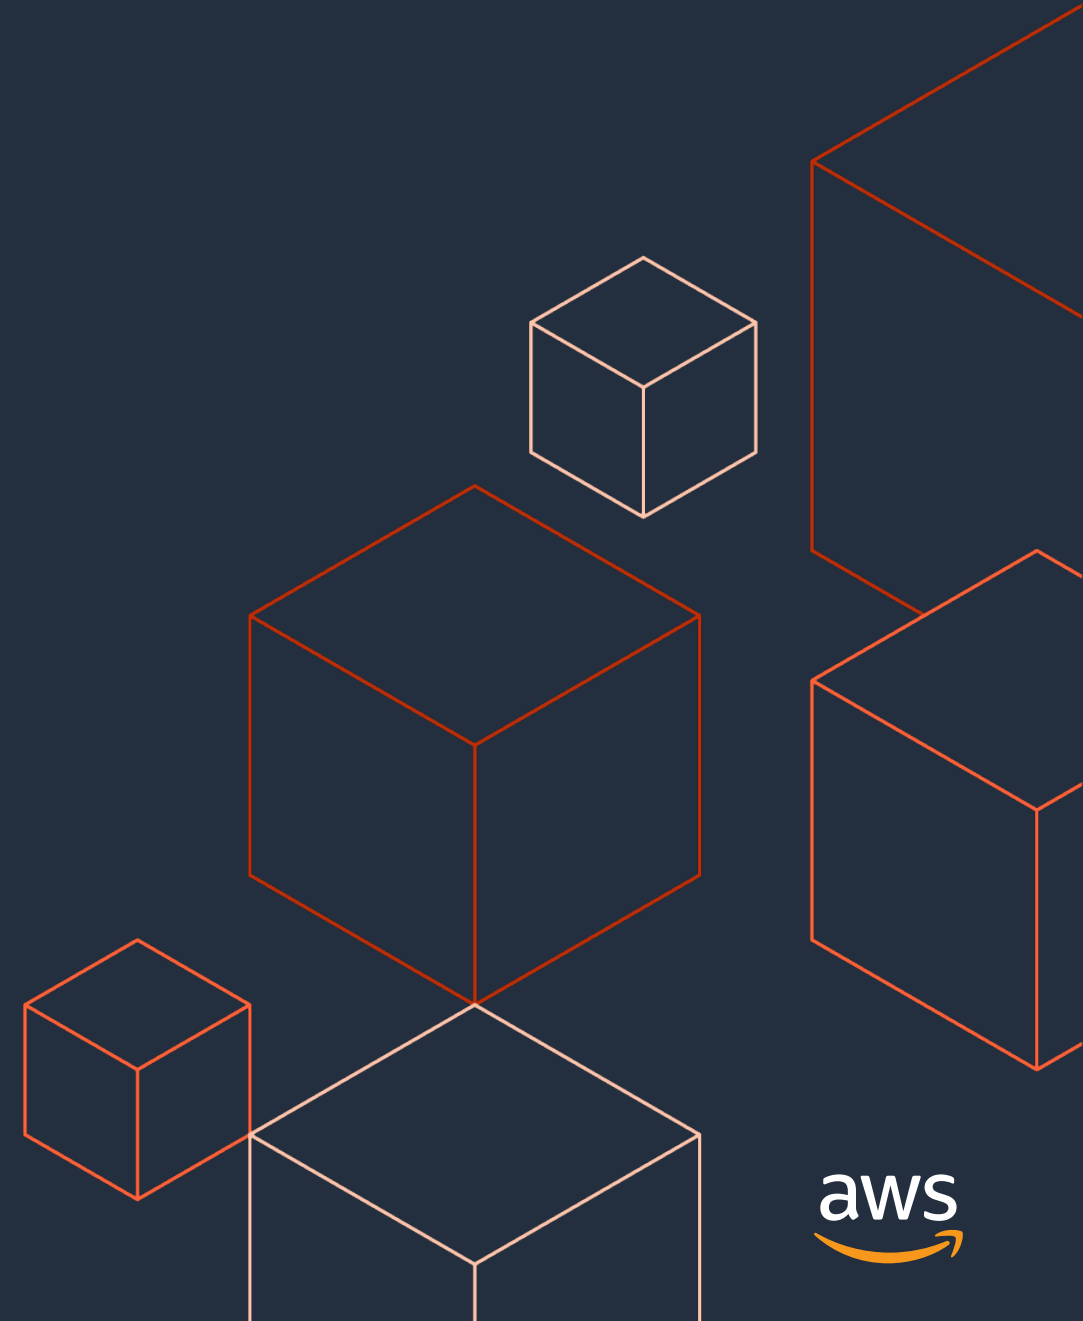
- HA Cluster setup and data replication for high availability
- Security upgrades will require downtime



High Operations

- Engineers/DBAs need to invest significant effort in provisioning fine tuning database configurations , patching and upgrades
- Deal with operational issues with database outages or availability drops

Modern Applications



Modern real-time applications require

Performance, Scale and Availability



E-Commerce



Media streaming



Social media



Online gaming

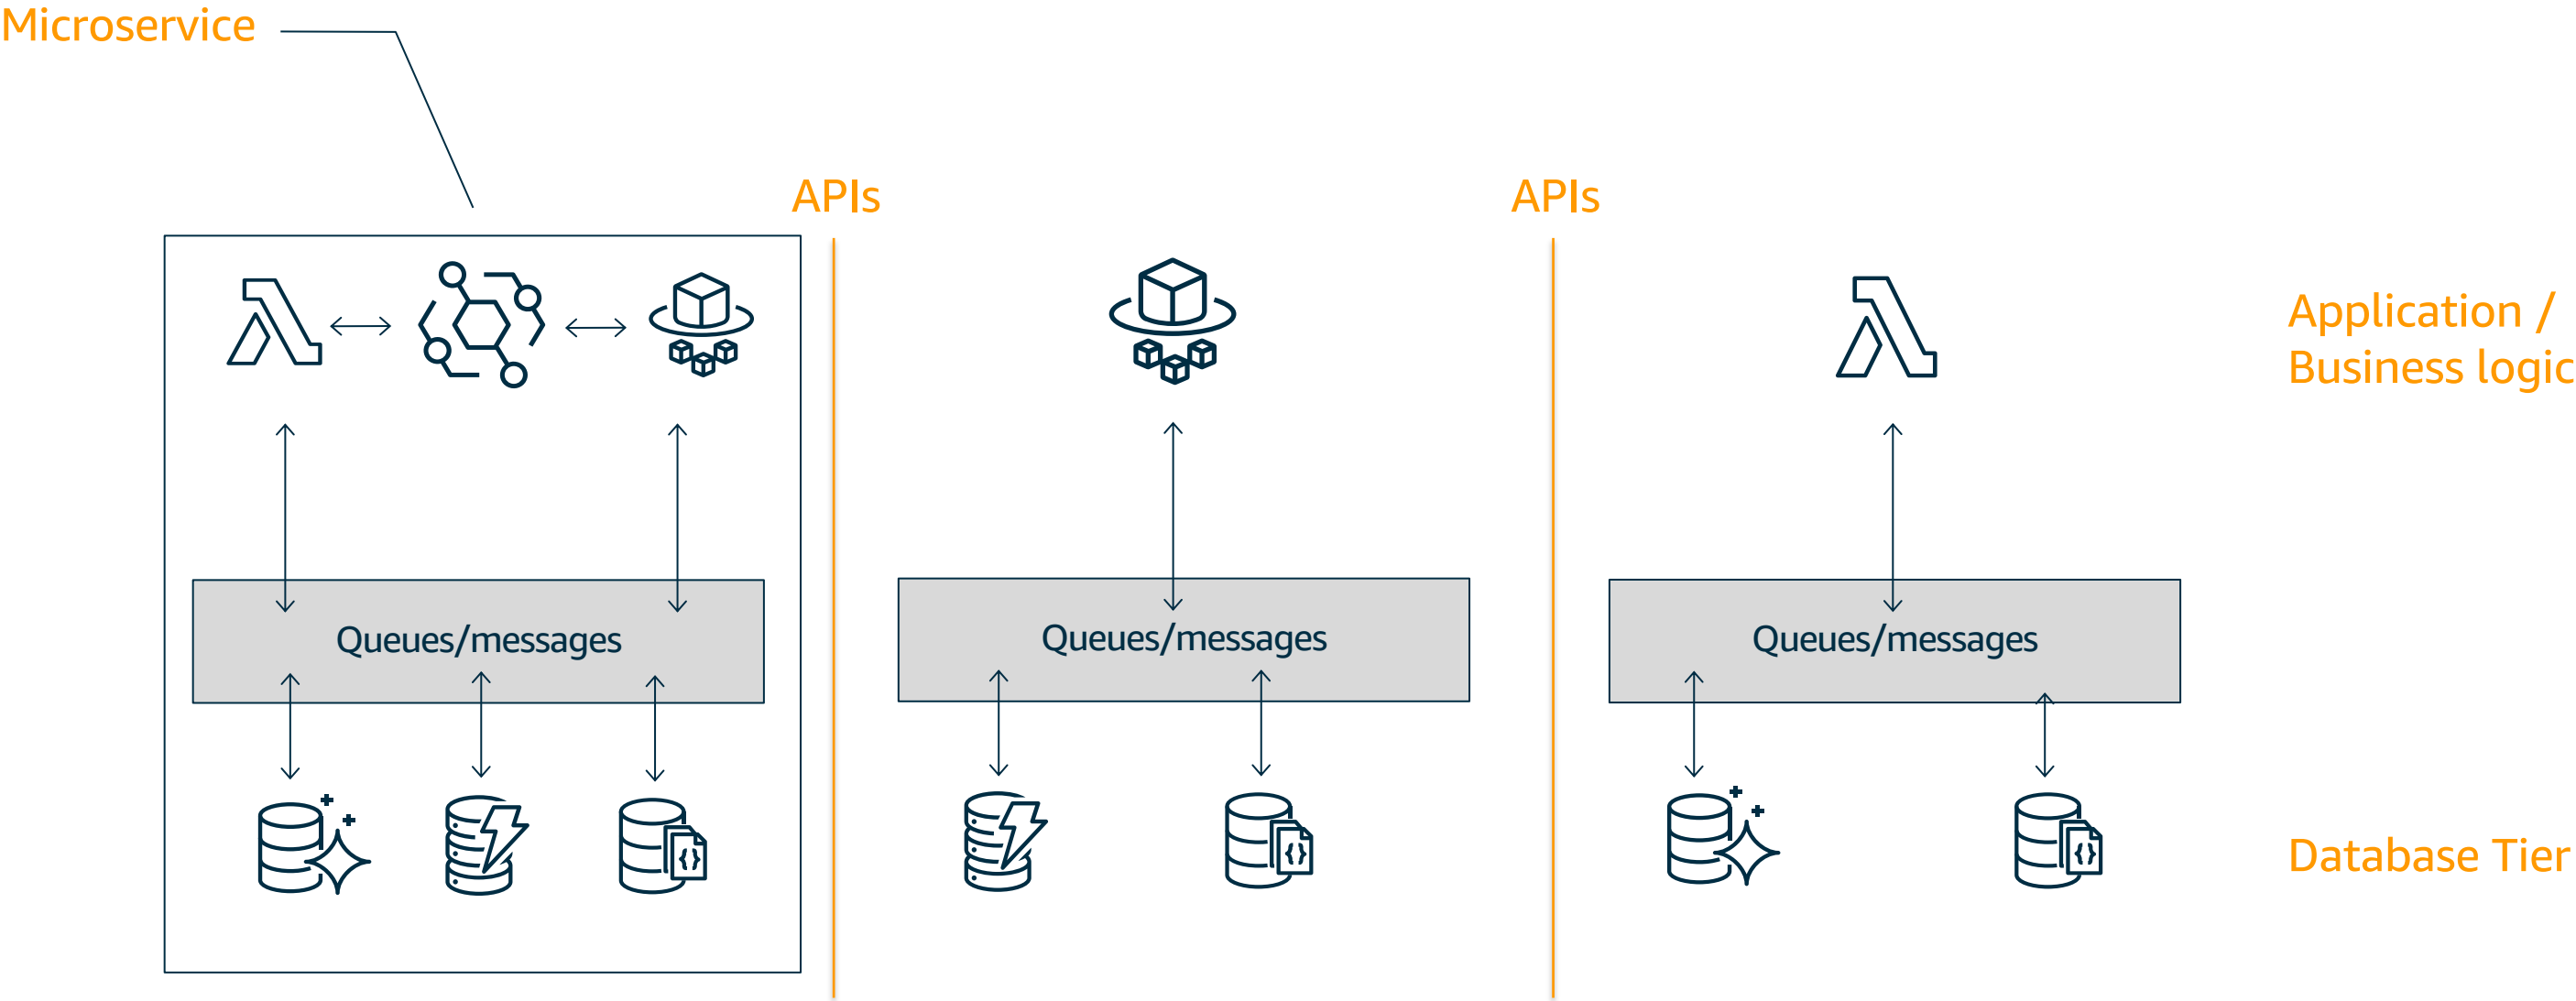


Shared economy

Users	1M+
Data volume	Terabytes—petabytes
Locality	Global
Performance	Microsecond latency
Request rate	Millions per second
Access	Mobile, IoT, devices
Scale	Up-down-out-in
Economics	Pay-as-you-go
Developer access	Open API

To keep up with requirements of modern applications...

customers are shifting to microservice architectures with purpose-built databases.



A one size fits all database doesn't fit anyone



Werner Vogels ✓

@Werner

Follow



A one size fits all database doesn't fit anyone
- #allthingsdistributed #aws



A one size fits all database doesn't fit anyone


The days of the one-size-fits-all monolithic database are behind us, and developers are using a multitude of purpose-built databases.

allthingsdistributed.com


Purpose Built Databases




Common data models and use cases




Relational




Key-value



Document




In-memory



Graph



Wide-column



Time-series



Ledger


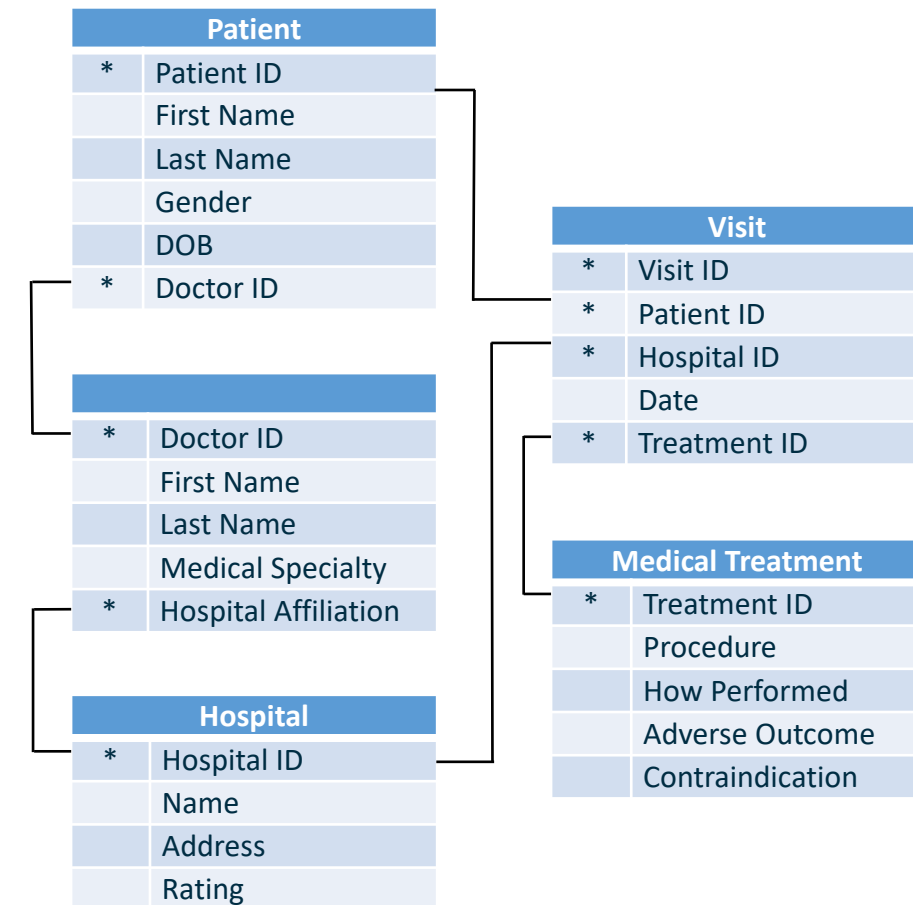
Purpose Built Databases: Use cases



Relational Use Case

- Divide data among tables
- Highly structured
- Relationships established via keys
- Data accuracy and consistency enforced by the system

```
SELECT
    d.first_name, d.last_name
FROM
    doctor as d,
    hospital as h
WHERE
    d.hospital = h.hospital_id
    AND h.name = Mercy
```



Amazon RDS

Managed relational database service with a choice of popular database engines

Amazon
Aurora


MySQL

PostgreSQL

MariaDB


Microsoft SQL Server

ORACLE




Easy to administer

No need to provision infrastructure, install, and maintain DB software



Available & durable

Automatic Multi-AZ data replication; automated backup, snapshots, and failover



Highly scalable

Scale DB compute and storage with a few clicks; minimal downtime for your application




Fast & secure

SSD storage and guaranteed provisioned I/O; data encryption at rest and in transit


Amazon Aurora

MySQL and PostgreSQL-compatible relational database built for the cloud




Performance & scalability

5x throughput of standard MySQL and 3x of standard PostgreSQL; scale-out up to 15 read replicas




Availability & durability

Fault-tolerant, self-healing storage; six copies of data across three AZs; continuous backup to S3



Highly secure


Network isolation, encryption at rest/transit



Fully managed

Managed by RDS: no server provisioning, software patching, setup, configuration, or backups

Key-value Use Case



- Simple key value pairs
- Partitioned by keys
- Consistent performance at scale


```
// Status of Hammer57
GET {
  TableName:"Gamers",
  Key: {
    "GamerTag":"Hammer57",
    "Type":"Status" } }

// Return all Hammer57
QUERY {
  TableName:"Gamers",
  KeyConditionExpression:"GamerTag = :a",
  ExpressionAttributeValues: {
    ":a":"Hammer57" } }
```

Gamers				
Primary Key		Attributes		
Gamer Tag	Type			
Hammer57	Rank	Level	Points	Tier
		87	4050	Elite
	Status	Health	Progress	
		90	30	
	Weapon	Class	Damage	Range
		Taser	87%	50
FluffyDuffy	Rank	Level	Points	Tier
		5	1072	Trainee
	Status	Health	Progress	
		37	8	

DynamoDB

Fast and flexible NoSQL database service for any scale



Performance at scale

- Handles millions of requests per second
- Delivers single-digit-millisecond latency
- Automated global replication
- Advanced streaming with Amazon Kinesis Data Streams for Amazon DynamoDB



No servers to manage

- Maintenance free
- Auto scaling
- On-demand capacity mode
- Change data capture for integration with AWS Lambda, Amazon Redshift, and Amazon OpenSearch Service



Enterprise ready

- ACID transactions
- Encryption at rest
- Continuous backups (supporting PITR), and on-demand backup and restore
- Export table data to Amazon S3
- PartiQL (a SQL-compatible query language) support

Wide-column Use Case




Amazon Keyspaces for Apache Cassandra

fully managed Apache Cassandra-compatible database service




Cassandra-compatible

Works with Apache 2.0-licensed
Cassandra drivers and developer
tools



Fast & scalable

Consistent, single-digit millisecond
read and write performance at any
scale



Fully managed (Serverless)


No servers to provision, patch, or
manage
No software to install, maintain, or
operate




Highly available & Secure

Data is replicated 3 times across
multiple AWS Availability Zones
Encryption at-rest

Document Use Case



Store, query, & index JSON
data natively



```
for i in response['businesses']:
    col.insert_one(i)
```


```
db.businesses.aggregate([ { $group: { _id "$price", ratingAvg: { $avg: "$rating" } } } ])
```

```
db.businesses.find({ $and: [{"price" : "$"}, {"rating": { $gt: 4.5 } } ]})
```

```
db.businesses.createIndex( { review_count: -1 } )
```


Amazon DocumentDB

Scale enterprise workloads with ease using a fully managed native JSON document database



Fast and scalable

- Decoupled compute and storage support independent scaling
- Scale to millions of reads using read replicas for instance-based clusters
- Scale to over a million of writes and reads and 1 PB of storage with Elastic Clusters



Fully managed

- Improve productivity and lower total cost of ownership by removing undifferentiated database management tasks



Enterprise ready


- High availability and durability
- Global Clusters provide local reads and disaster recovery from Region-wide outages
- Built-in security best practices with encryption-in-transit and encryption-at-rest, Amazon VPC, and AWS KMS



MongoDB compatible

- Supports hundreds of APIs, operators, and data types
- MongoDB APIs, drivers, and tools can be used with little to no change

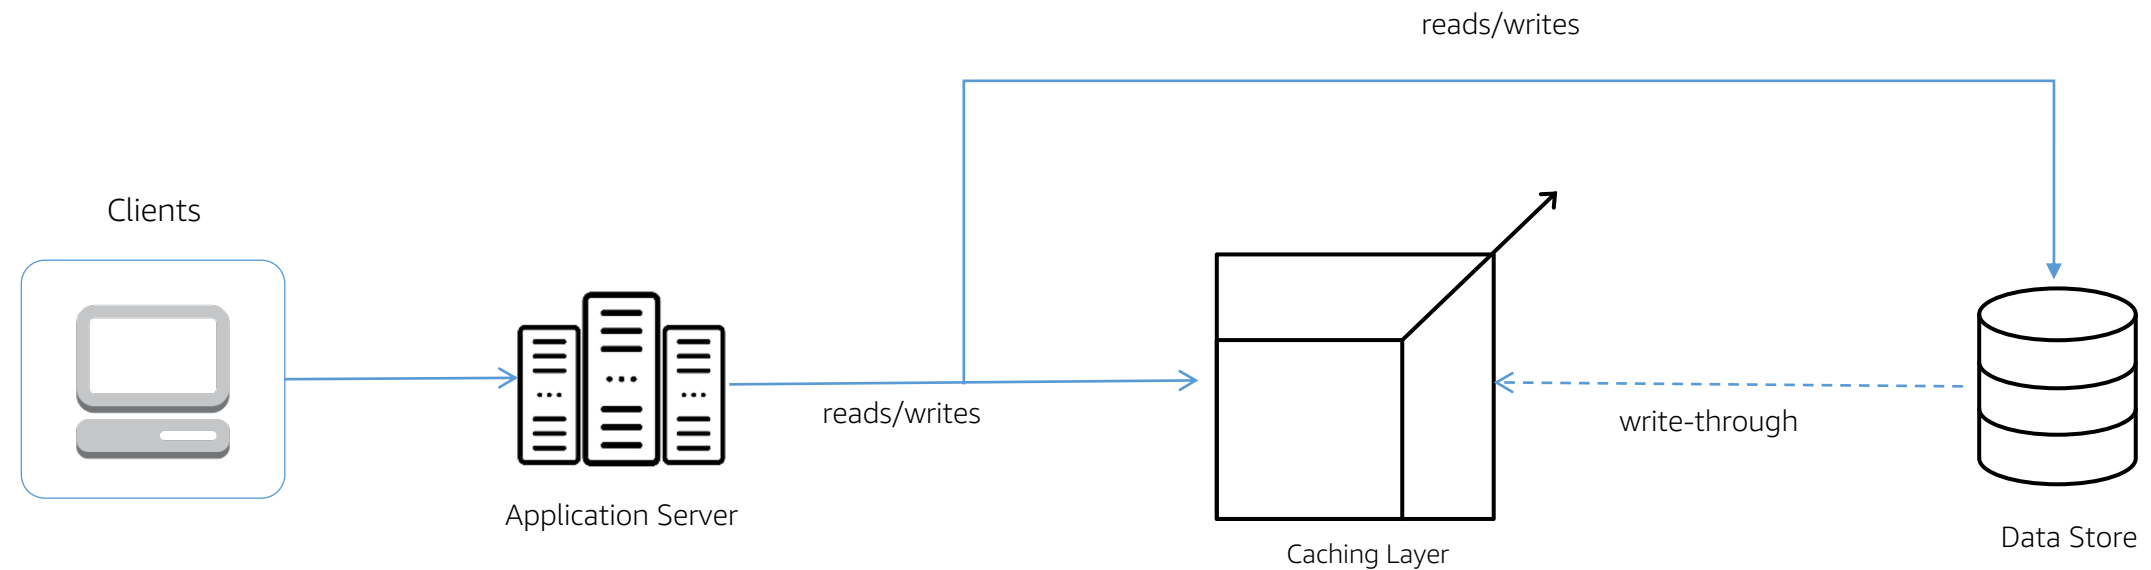
In-memory Use Case



Latency: Extremely low, microsecond to millisecond responses


Request Rate: Thousands to millions of reads and/or writes per second

Data Volume: Will scale up to over 100TB




Amazon ElastiCache

Fully managed, Redis or Memcached compatible, low-latency, in-memory data store




Extreme performance

In-memory data store and cache
for sub-millisecond response times



Fully managed

AWS manages all hardware
and software setup,
configuration, and monitoring



Easily scalable

Read scaling with replicas
Write and memory scaling
with sharding
Nondisruptive scaling

Amazon MemoryDB


Redis-compatible, durable, in-memory database service for ultra-fast performance

Redis
compatible



Fully compatible with
open source Redis
Cluster

Extreme
performance




In-memory database
for microsecond reads

Secure
and reliable




Network isolation, encryption
at rest/transit, multi AZ, and
automatic failover

Easily scales to
massive workloads




Scale writes and
reads with sharding
and replicas

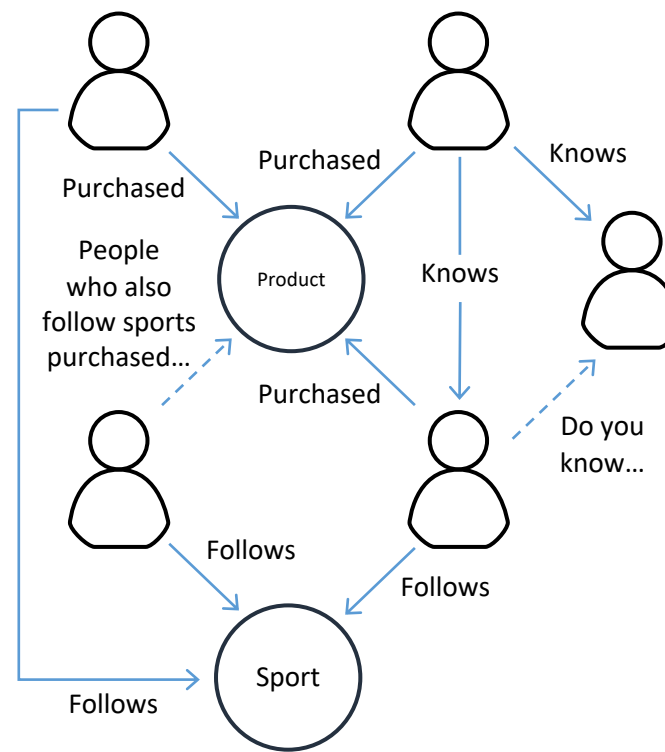
Graph Database Use Case



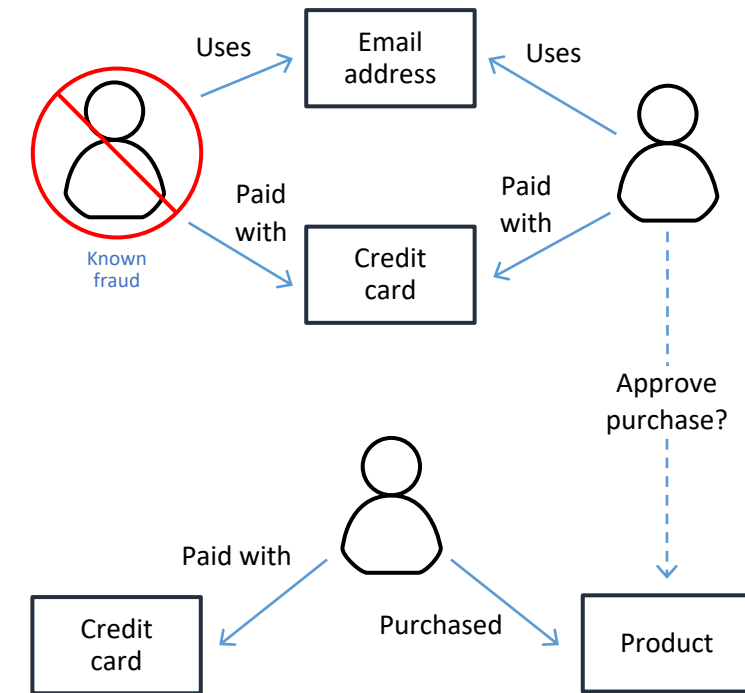
Social news feed



Recommendations



Retail fraud detection



Amazon Neptune

Fast, reliable graph database built for the cloud

Open



Supports Apache TinkerPop &
W3C RDF graph models

Fast



Query billions of relationships
with millisecond latency

Reliable





6 replicas of your data across 3
AZs with full backup and restore

Easy




Build powerful queries easily
with Gremlin and SPARQL

Time-series



Amazon Timestream

serverless time series database service for IoT and operational applications




1,000x faster at 1/10th the cost of relational databases

Collect fast moving time-series data from multiple sources at the rate of millions of inserts per second




Trillions of daily events

Capable of processing trillions of events daily; the adaptive query processing engine maintains steady, predictable performance



Analytics optimized for time series data

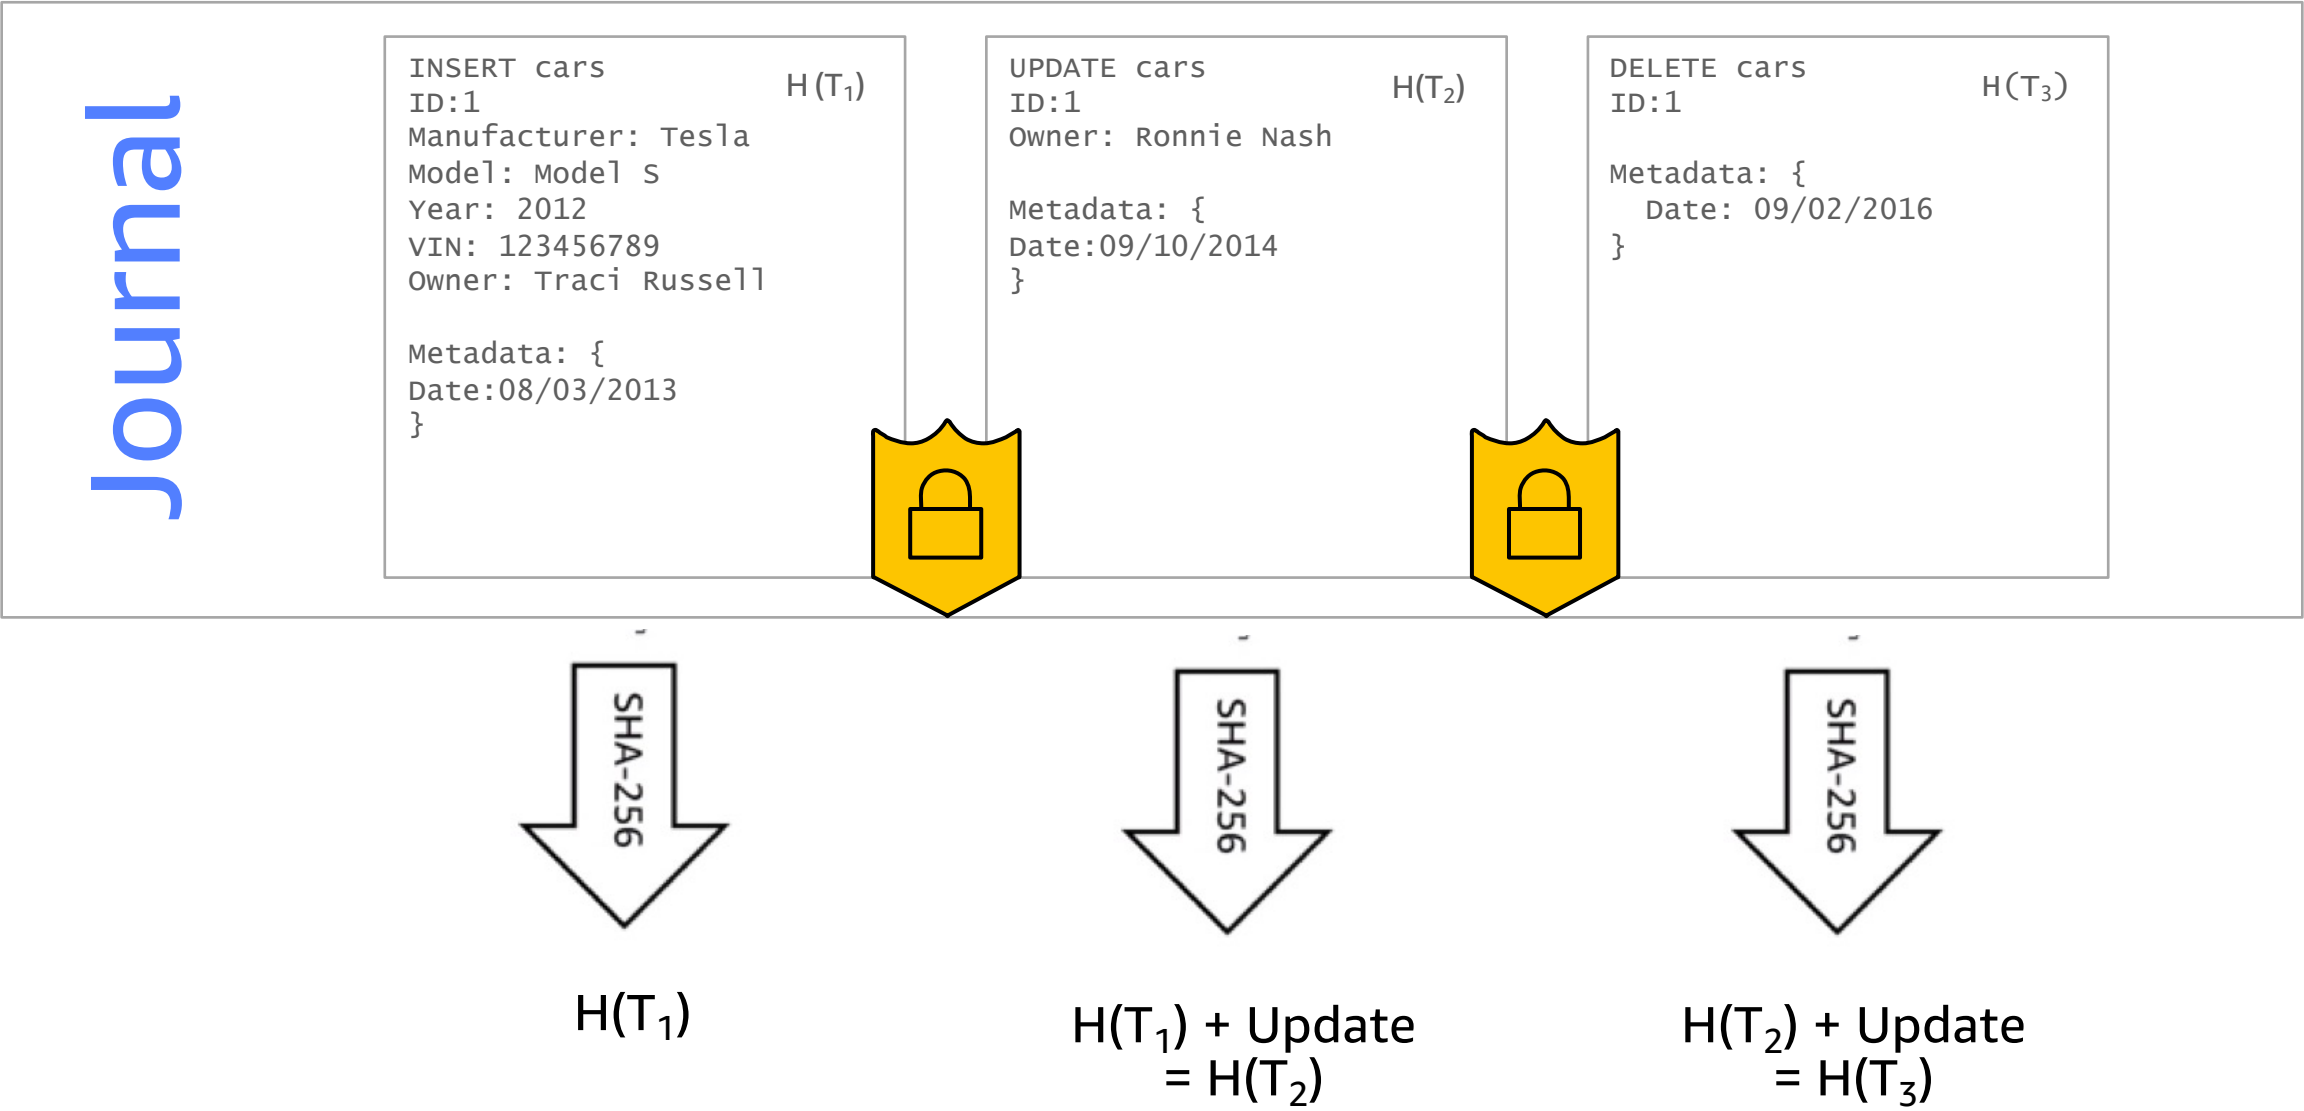
Built-in analytics for interpolation, smoothing, and approximation to identify trends, patterns, and anomalies



Serverless


No servers to manage; time-consuming tasks such as hardware provisioning, software patching, setup, & configuration done for you

Ledger Use Case




Amazon QLDB

fully managed ledger database




Immutable

Maintains a sequenced record of all changes to your data, which cannot be deleted or modified; you have the ability to query and analyze the full history




Cryptographically verifiable

Uses cryptography to generate a secure output file of your data's history



Highly scalable

Executes 2–3X as many transactions as ledgers in common blockchain frameworks




Easy to use

Easy to use, letting you use familiar database capabilities like SQL APIs for querying the data











Best Practices with purpose built databases

- Structure of your data
 - Data types
 - Flat vs Json
 - Volume and size of data over time
- Data Modeling
 - Choosing the key for uniqueness and join across data stores
- Access Pattern
 - Key vs Range lookup vs Full search
 - Ad-Hoc/Ops Queries
 - Access pattern changes over time?
- Data consistency and durability
 - Strong vs Eventual
 - Source of truth
 - Data validation / reconciliation
- Performance
 - Concurrency
 - Transaction throughput/response time (p50/p99)
- Data Integration
 - Federated Query
 - ETL
 - Data Lake

Summary



Summary

	Relational	Key-value	Document	In-memory	Graph	Time-series	Ledger	Wide Column
Data Models	Referential integrity, ACID transactions, schema-on-write	High throughput, low-latency reads and writes, endless scale	Store documents and quickly access querying on any attribute	Query by key with microsecond latency	Quickly and easily create and navigate relationships between data	Collect, store, and process data sequenced by time	Complete, immutable, and verifiable history of all changes to application data	Scalable, highly available, and managed Apache Cassandra-compatible service
Common Use Cases	Lift and shift, ERP, CRM, finance	Real-time bidding, shopping cart, social, product catalog, customer preferences	Content management, personalization, mobile	Leaderboards, real-time analytics, caching	Fraud detection, social networking, recommendation engine	IoT applications, event tracking	Systems of record, supply chain, health care, registrations, financial	Build low-latency applications, leverage open source, migrate Cassandra to the cloud
	<div> Amazon Aurora</div> <div> Amazon Relational Database Service (RDS)</div>	<div> Amazon DynamoDB</div>	<div> Amazon DocumentDB</div>	<div> Amazon ElastiCache</div> <div> Amazon MemoryDB</div>	<div> Amazon Neptune</div>	<div> Amazon Timestream</div>	<div> Amazon QLDB</div>	<div> Amazon Keyspaces (for Apache Cassandra)</div>



Airbnb uses different databases based on the purpose

User search history: **Amazon DynamoDB**

- Massive data volume
- Need quick lookups for personalized search

Session state: **Amazon ElastiCache**

- In-memory store for submillisecond site rendering

Relational data: **Amazon RDS**

- Referential integrity
- Primary transactional database