



# Amazon API Gateway

## Overview

Eren Akbaba

Solutions Architect, AWS  
eakbaba@amazon.com

# Agenda

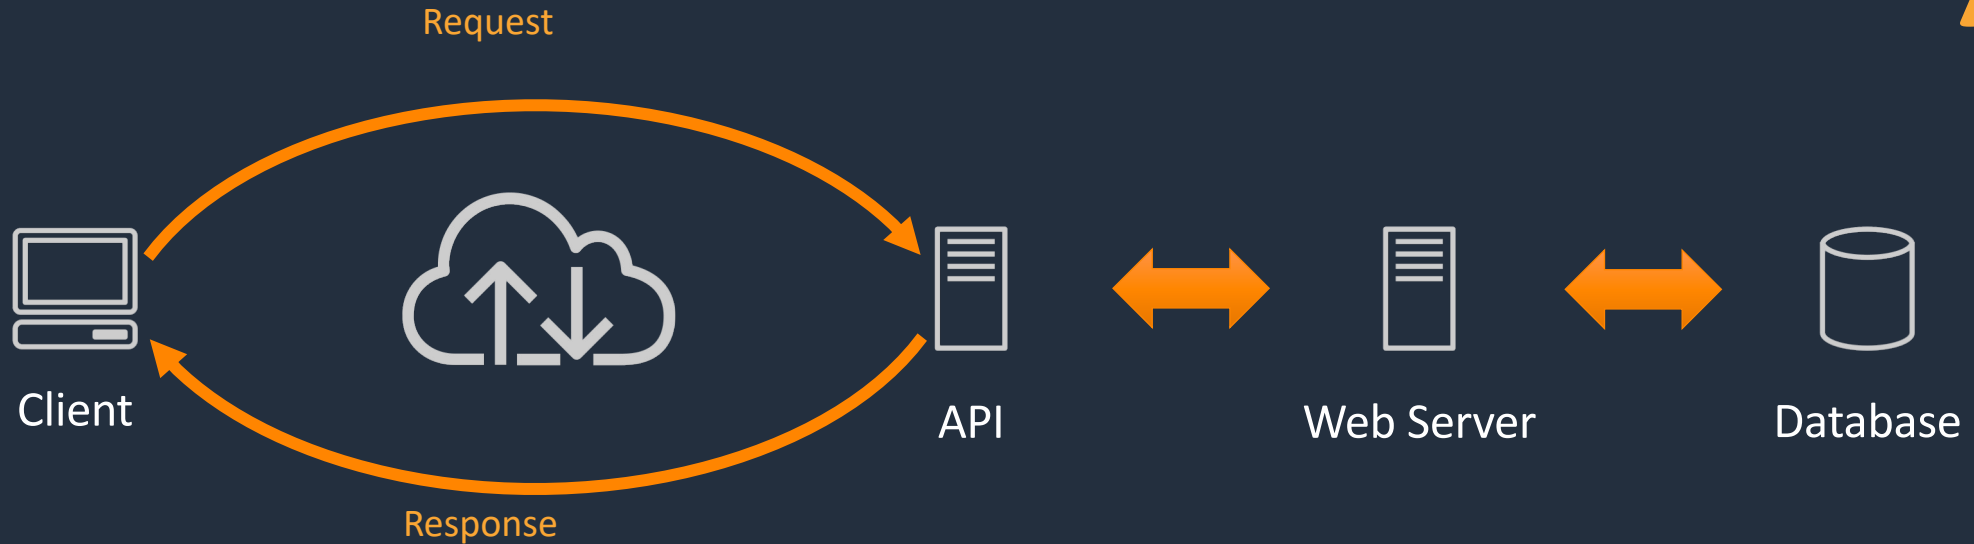
- What is API Gateway
- API Types
- Integrations
- Protecting APIs
- Validation and Transformation
- Stages and Versioning
- Custom Domain Names
- Observability
- Other Features
- Pricing
- Best Practices

# What is API Gateway?



# Application Programming Interface (API)

“ In building applications, an API simplifies programming by abstracting the underlying implementation and only exposing objects or actions the developer needs. ”



Web-based companies and services offer APIs for developers to use, such as:

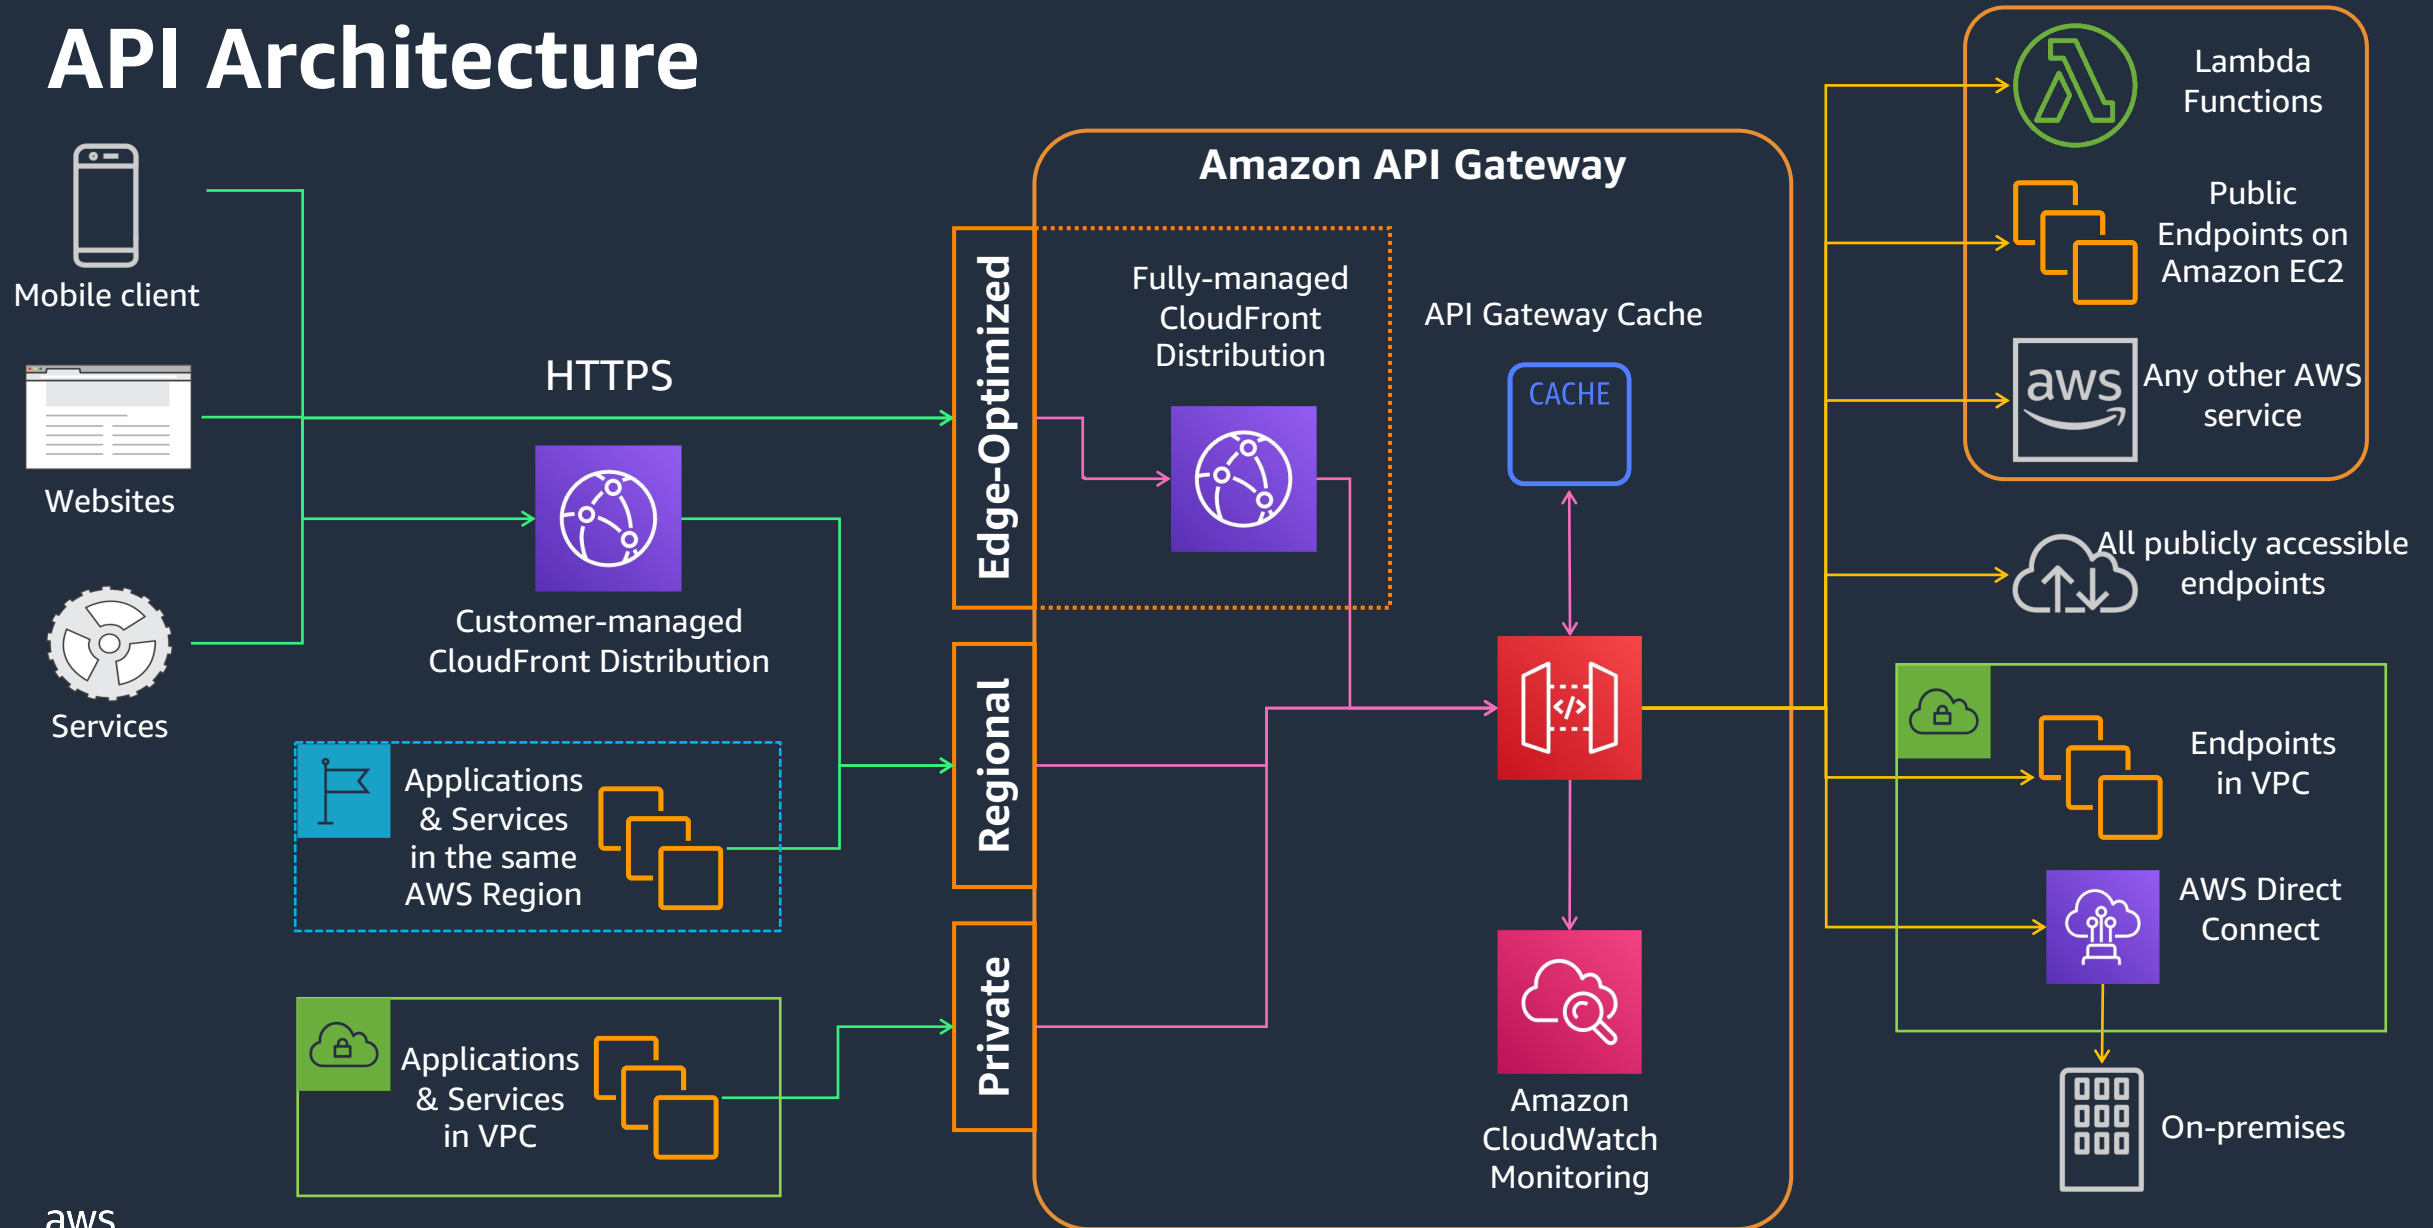
- Social Networks – Facebook, Twitter, etc
- Payment Processing – Amazon Pay, PayPal, etc

# Amazon API Gateway

Amazon API Gateway is a fully managed (serverless) service that makes it easy for developers to create, publish, maintain, monitor, and secure APIs at any scale.



# API Architecture



# API Types




# Supported Protocols

## RESTful APIs



- Request / Response
- HTTP Methods like GET, POST, etc
- Short-lived communication
- Stateless

## WebSocket APIs



- Serverless WebSocket
- 2 way communication channel
- Long-lived communication
- Stateful



# RESTful APIs

- Two flavors: REST API (v1) and HTTP API (v2)
- REST API is more feature rich, feature parity in HTTP API will take some time.
- HTTP API is built from the ground up:
  - Faster – up to 60% faster
  - Lower cost – up to 71% less expensive
  - Easier to use

# Endpoint types

## REST Edge-Optimized

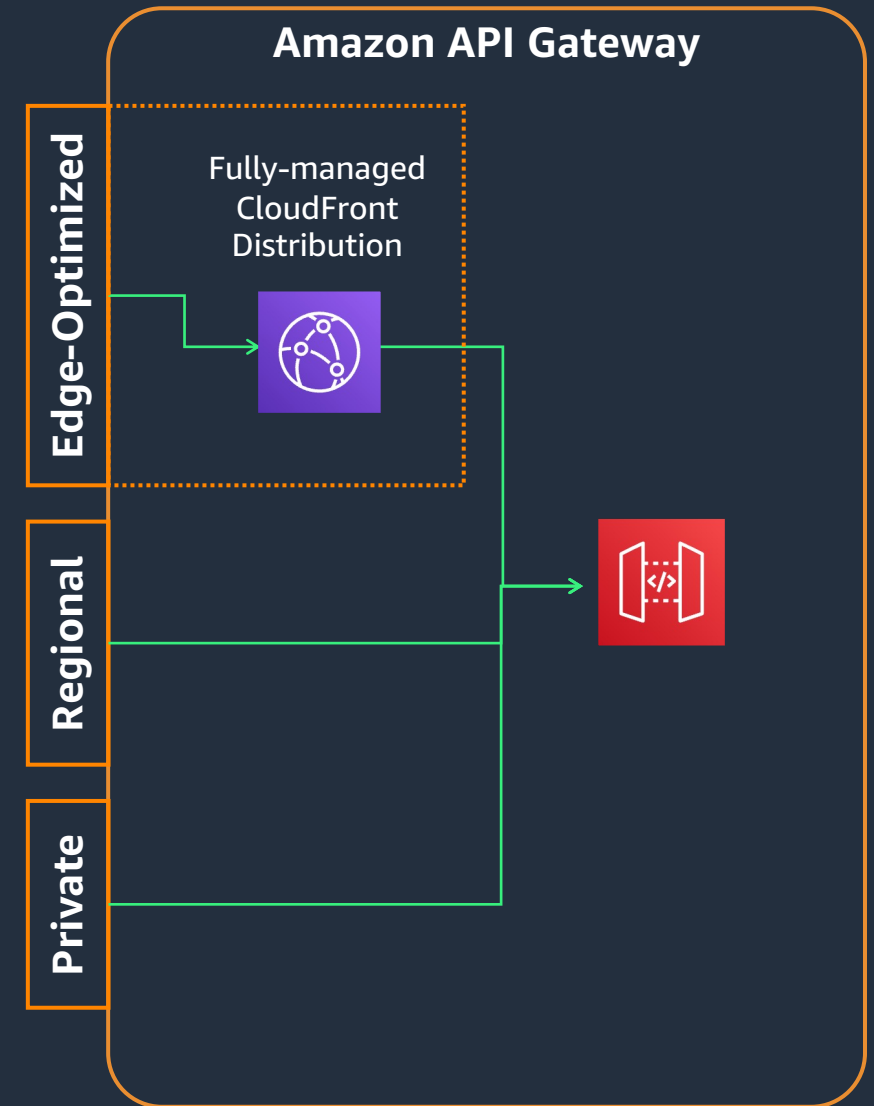
- Utilizes CloudFront to reduce TLS connection overhead (reduces roundtrip time)
- Designed for a globally distributed set of clients

## Regional

- Recommended API type for general use cases
- Designed for building APIs for clients in the same region

## REST Private

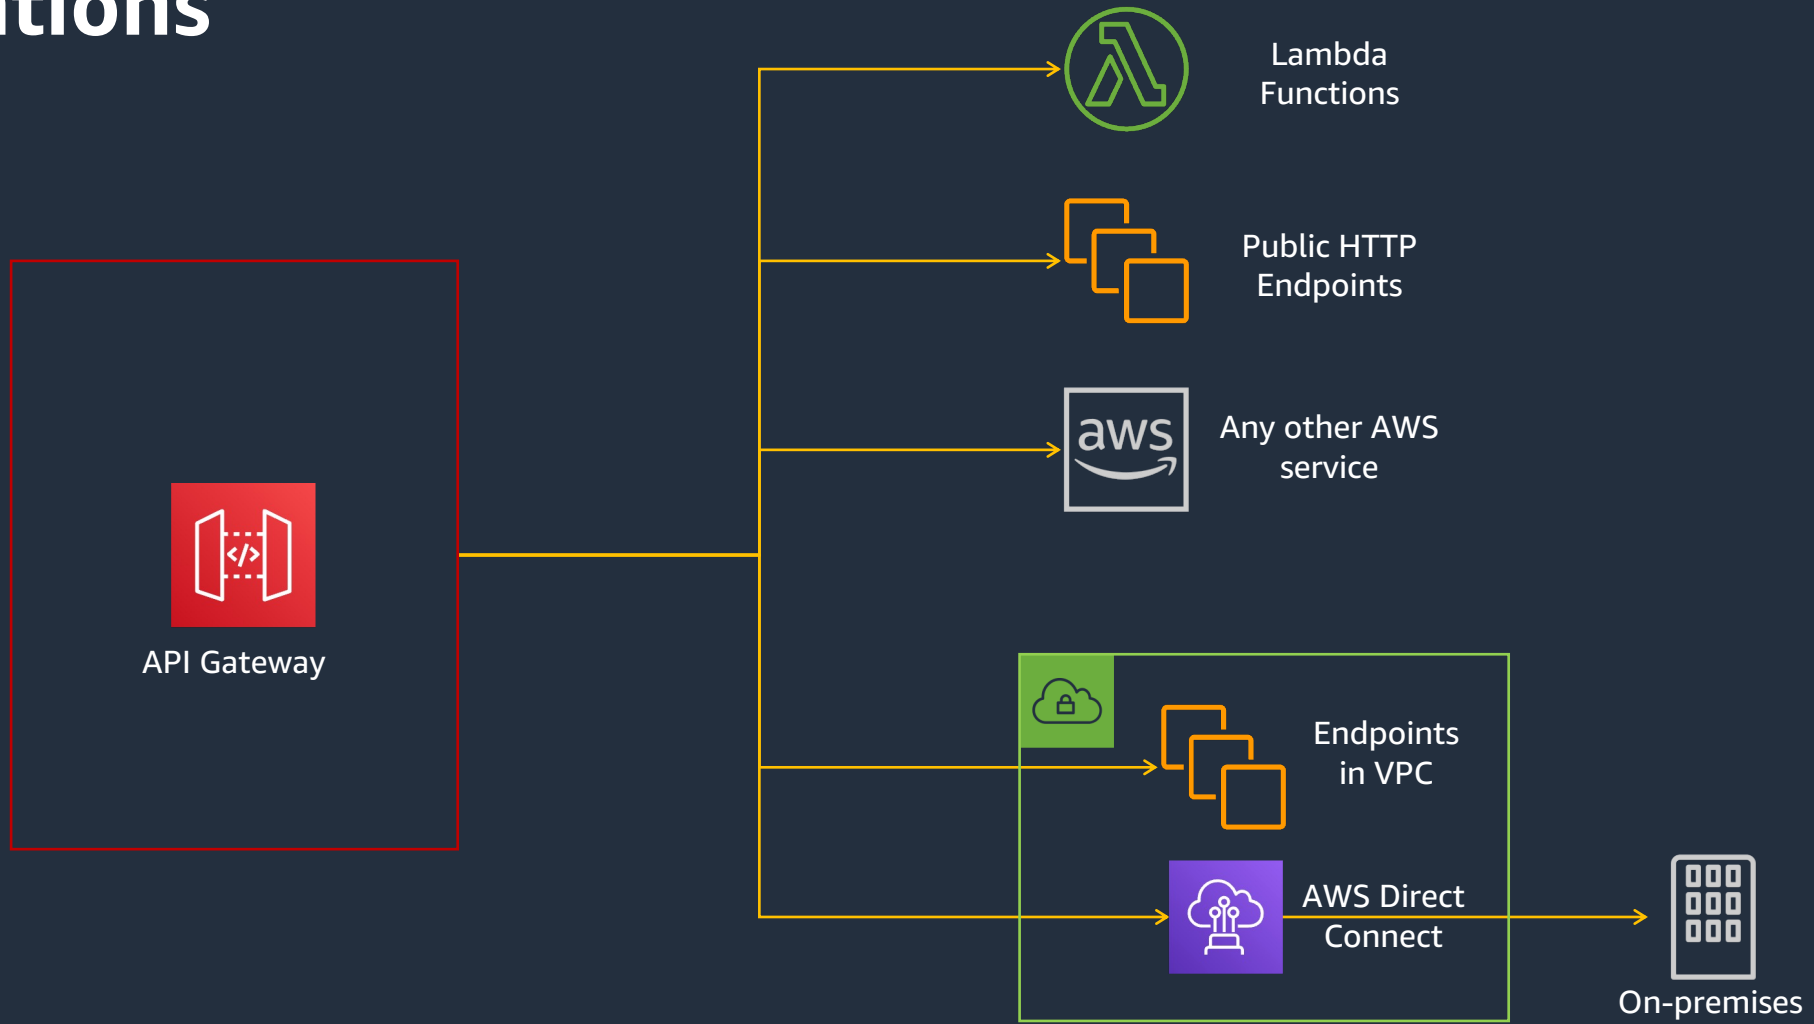
- Only accessible from within VPC (and networks connected to VPC)
- Designed for building APIs used internally or by private microservices



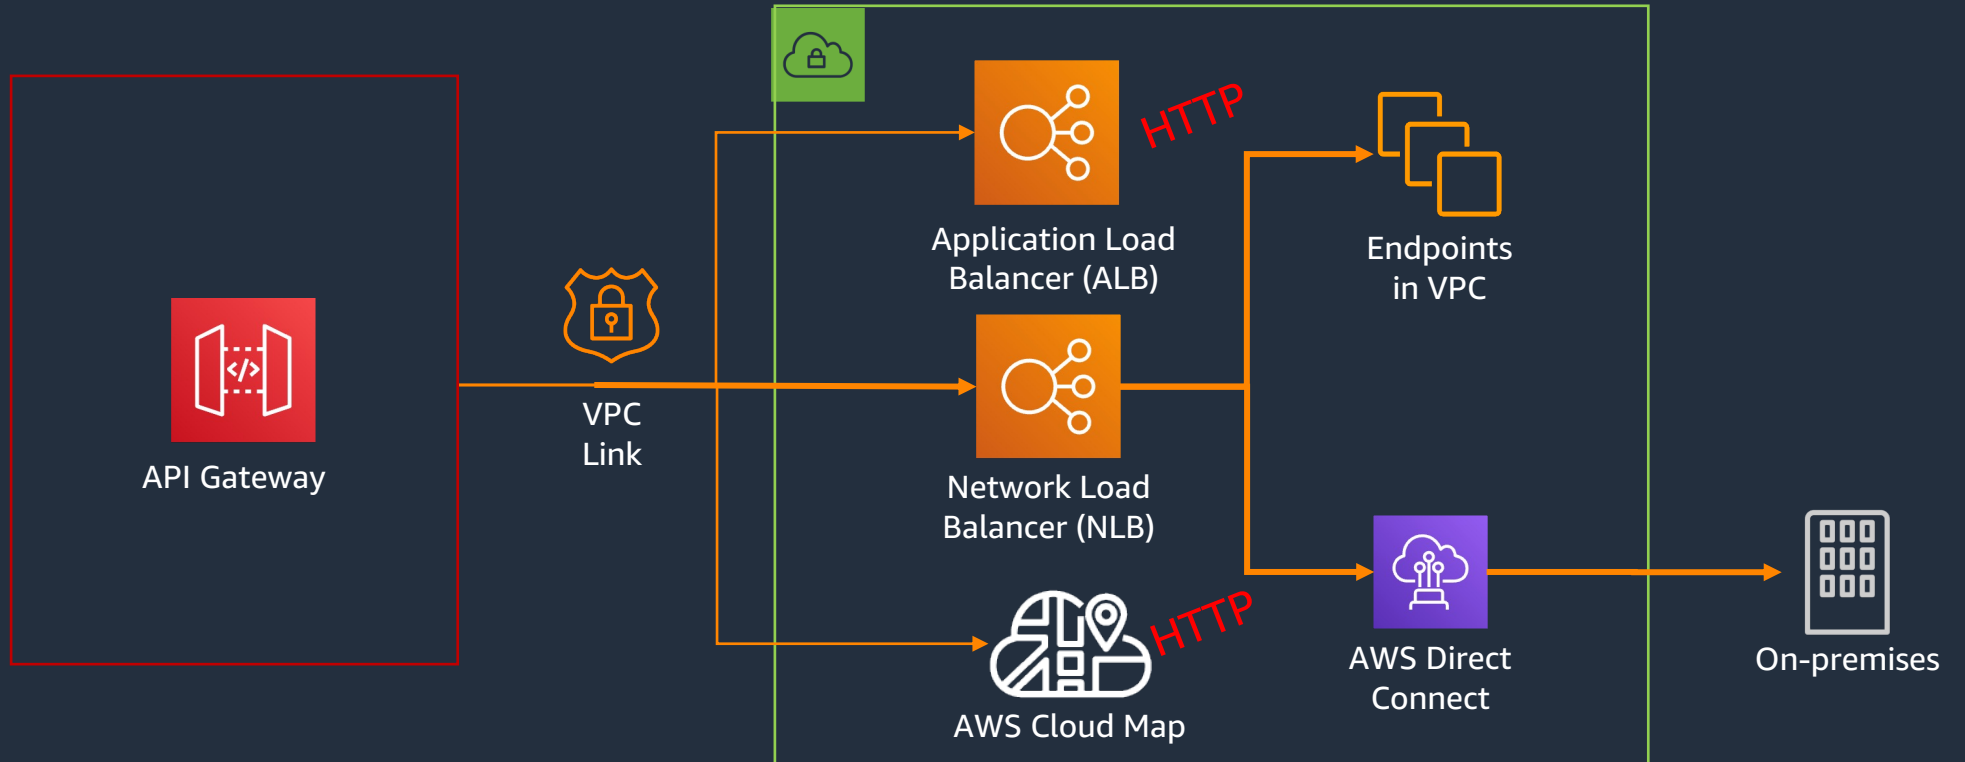
# Integrations




# Integrations




# Private Integrations – VPC Link



# The request cycle




# The request cycle



Only appears when an AWS WAF web access control list (ACL) is configured for enhanced security. During this phase, AWS WAF rules are evaluated and a decision is made on whether to continue or cancel the request.

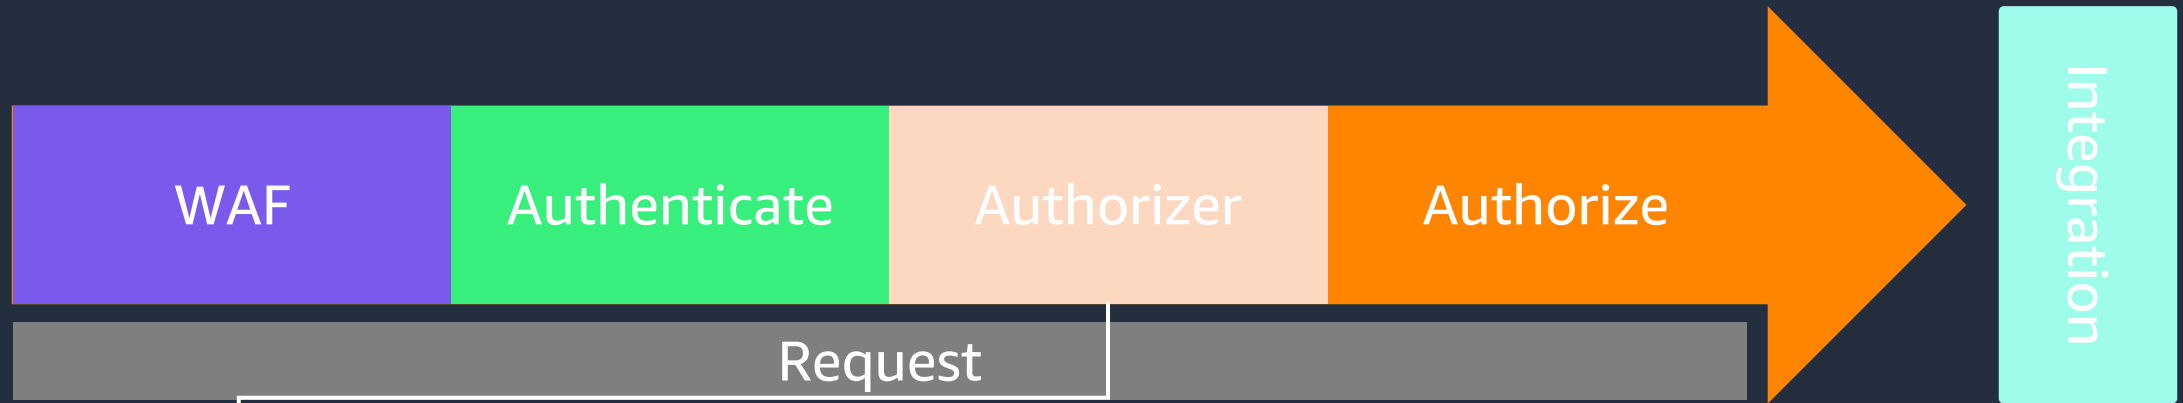
# The request cycle



Only present when AWS Identity and Access Management (IAM) authorizers are used. During this phase, the credentials of the signed request are verified. Access is granted or denied based on the client's right to assume the access role.

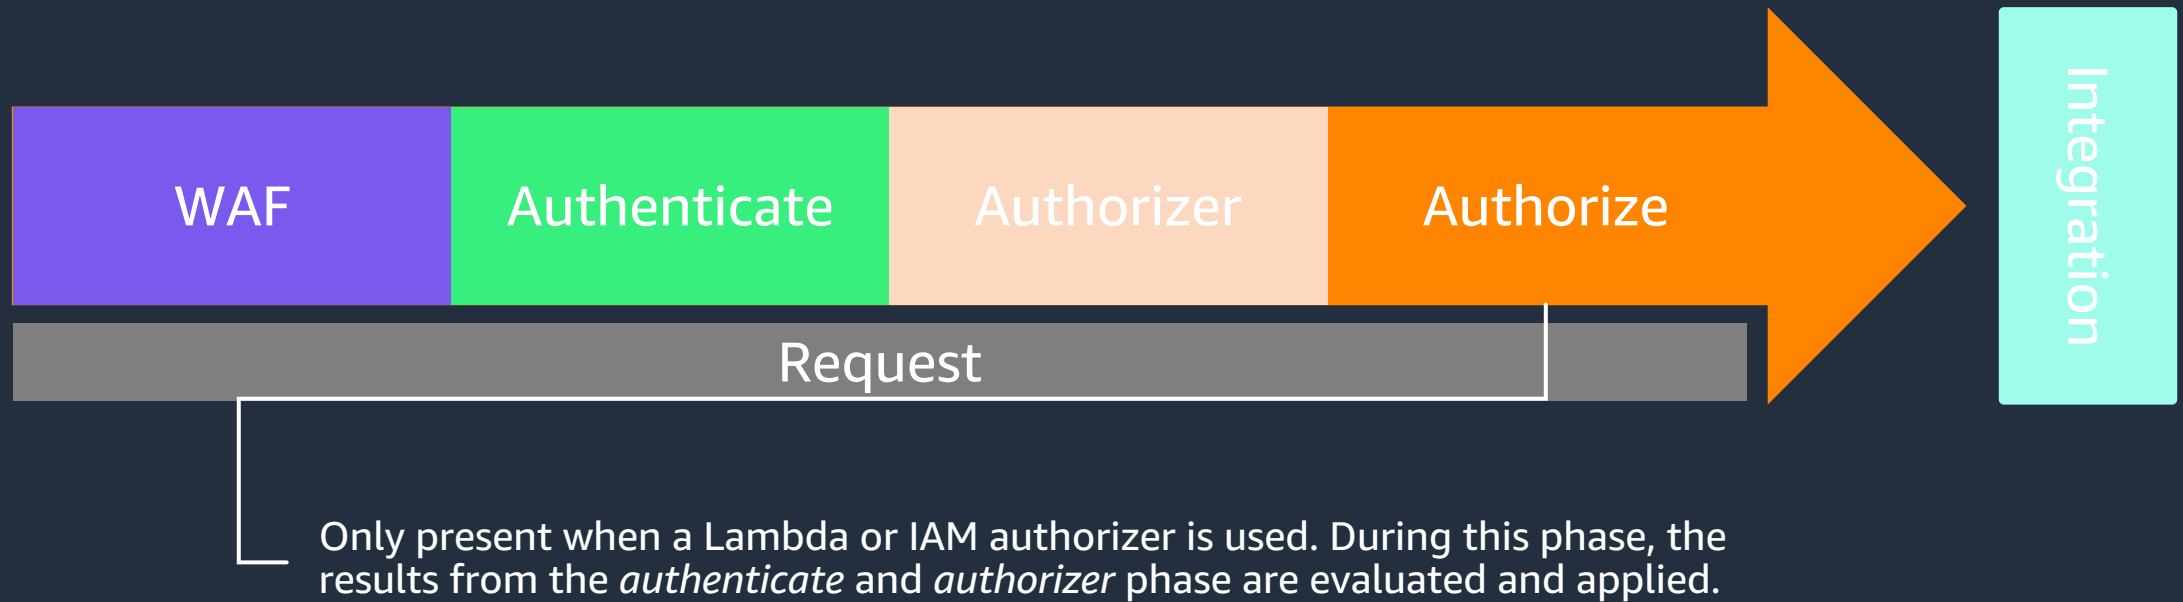


# The request cycle




Only present when a Lambda, JWT, or Amazon Cognito authorizer is used.  
During this phase, the authorizer logic is processed to verify the user's right to access the resource.

# The request cycle



REST

# Client Certificates



- Generate client-side SSL certificate using the API Gateway
- Allow backend to verify request coming from API Gateway using public key
- Expires after 365 days

# Protecting APIs



# API Security with Amazon API Gateway



# Types of authorization

**REST** ➤ **Resource Policy**

➤ **Mutual TLS**

**REST** ➤ **WAF**

**REST** ➤ **Cognito User Pools**

**HTTP** ➤ **JWT**

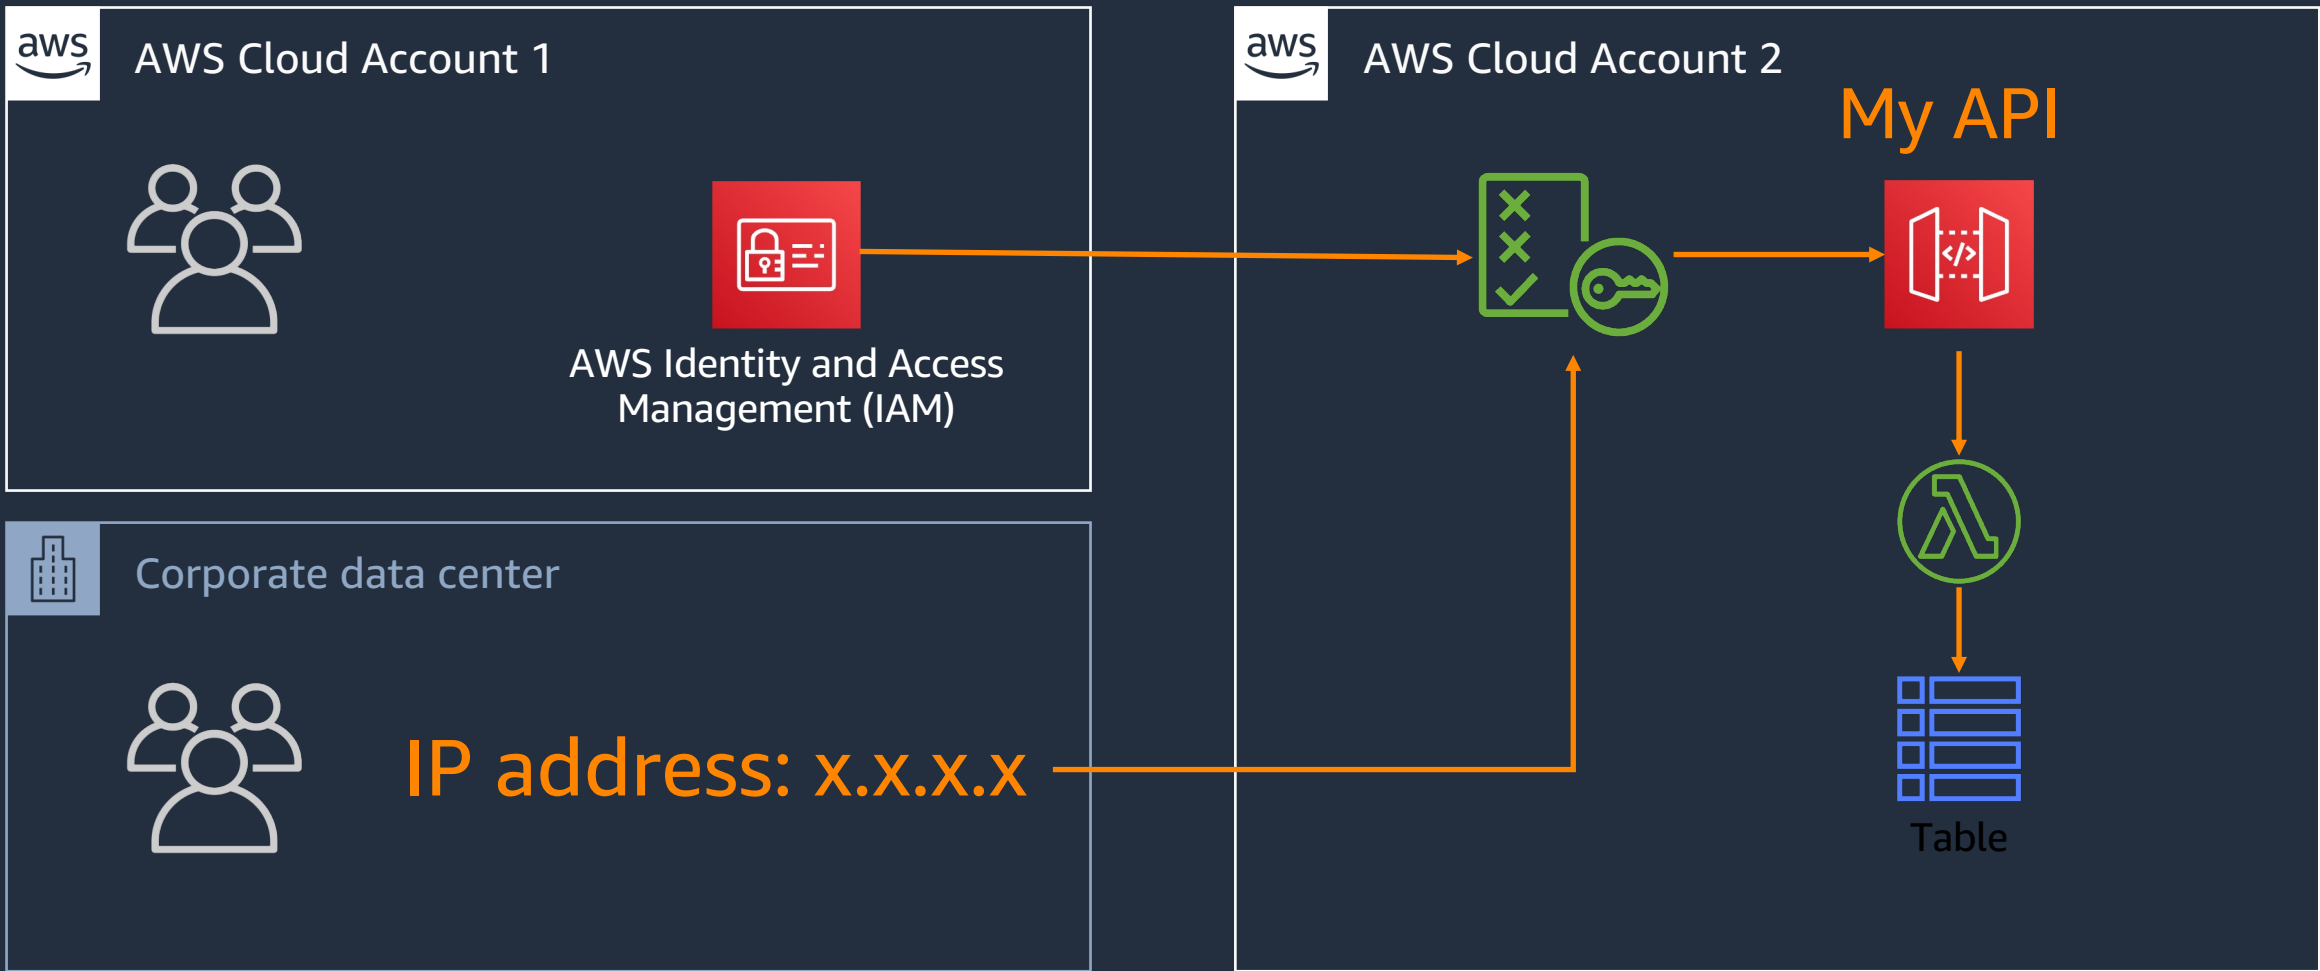
**+**

➤ **IAM**

➤ **Lambda Authorizer**

REST

# Resource Policies



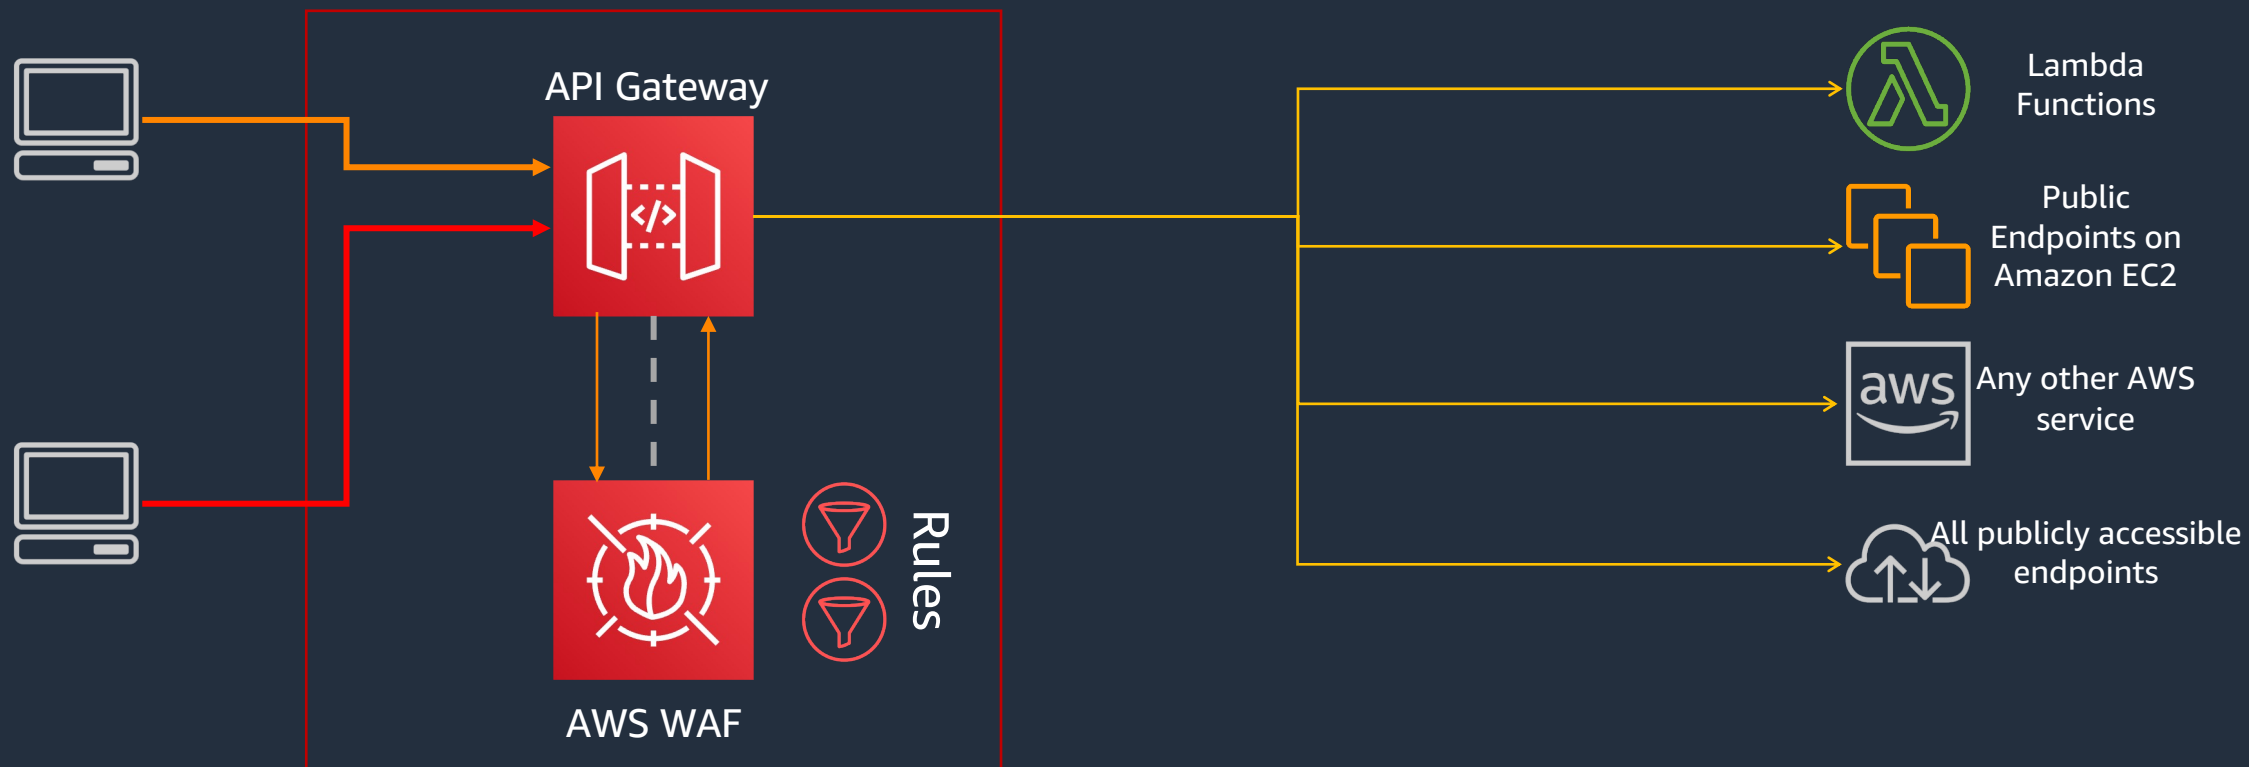
# Mutual TLS

- Requires two-way authentication between client and server
- Client must present X.509 certificates to verify identity
- Often used in Internet of Things (IoT) and business to business applications
- Only supported on Custom Domains





# REST Web Application Firewall (WAF)




# HTTP REST JWT/Cognito authorizer

- OAuth2 compliant (part of OpenID Connect - OIDC)
- Allows or denies access based on token validity and optional scopes
- Any required scopes for the route are validated in the token



# IAM authorizer


- Clients must use Signature Version 4 to sign their requests with AWS credentials
- Authorization token is decoded
- User is verified against the Identity & Access Management (IAM) service
- User must have *execute-api* access on the route to proceed



# Lambda authorizer

- Your custom logic to validate the request
- 2 payload options
  - Payload 1: must return an IAM policy that allows or denies access to your API route
  - Payload 2: can return IAM policy or Boolean
- Authorization can be cached

HTTP




# Throttling and usage plans <sup>REST</sup>

- Protect backend systems
- Prevents one customer from consuming all your backend system's capacity
- Let's you decide how to allocate capacity among your API consumers with quotas and request rates.

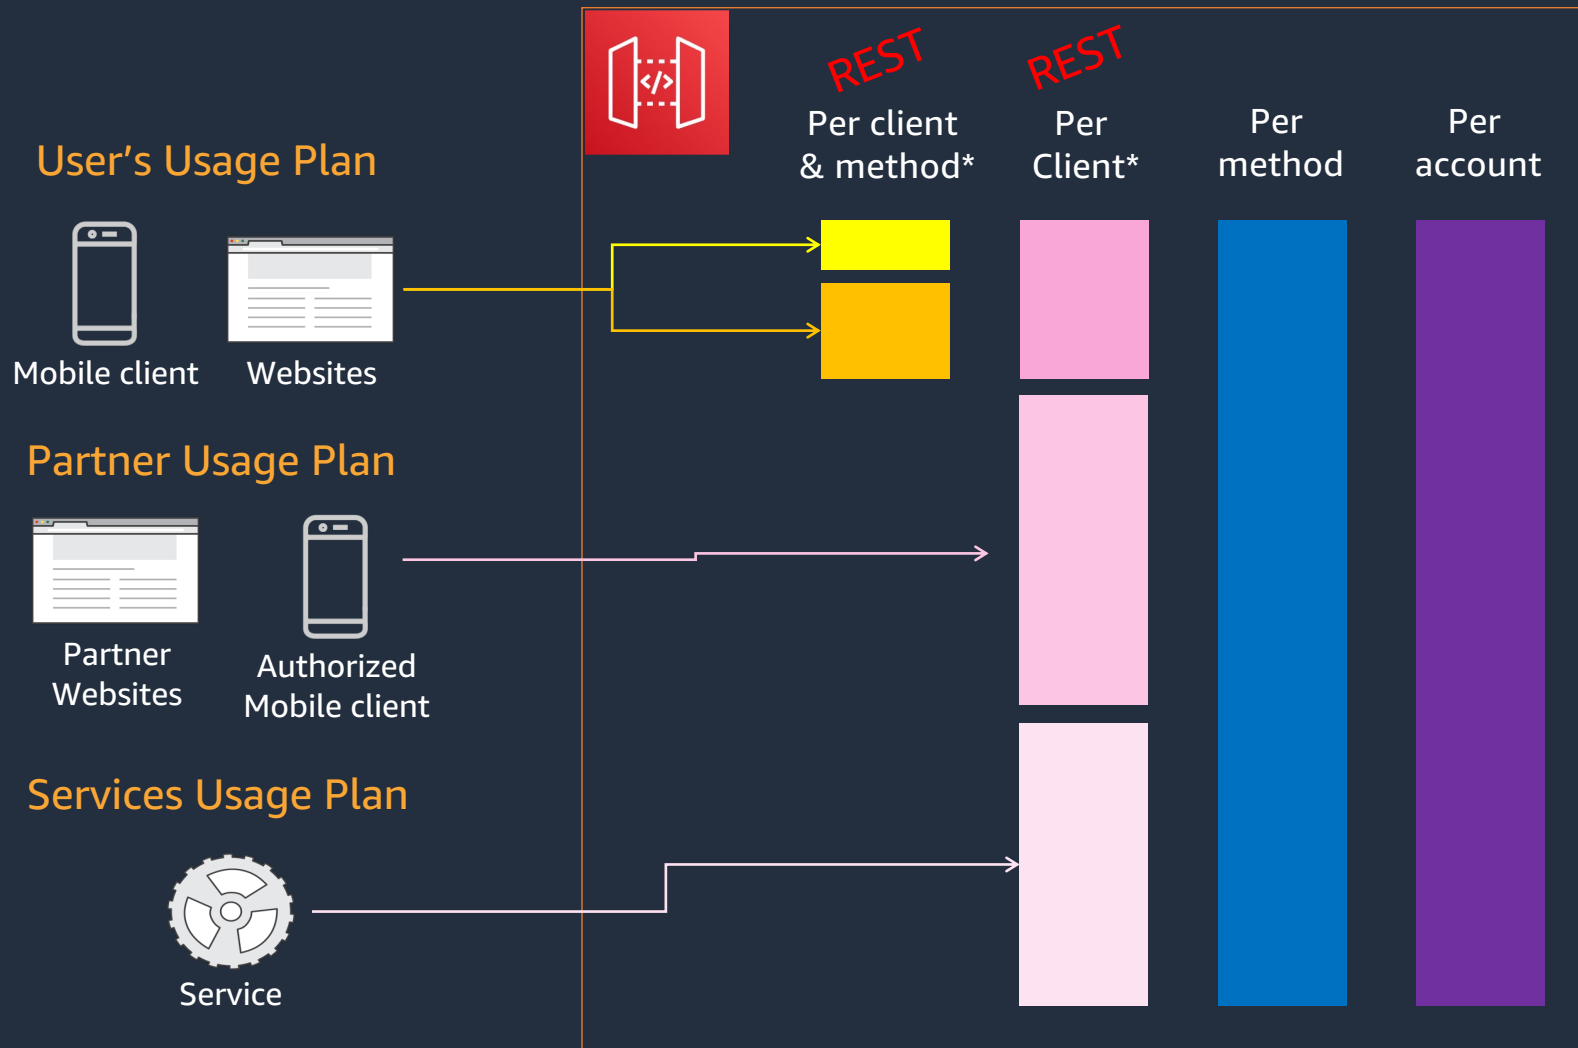
Professional plan users:  
10 RPS, up to 100 calls / day



Enterprise plan users:  
500 RPS, no limit on calls / day



# Four levels of Throttling



REST

\* Requires API Keys and Usage Plans



# REST Usage Plans and API Keys

## API Keys

- Alphanumeric string values that you distribute to clients (per user/client)
- Generated by API Gateway or you can imported from a CSV file
- Use API keys together with Usage Plans or Lambda authorizers to control access to your APIs

## Usage Plans

- Specifies who can access API stages and methods
- How much and how fast they can access the resource
  - Rate limit
  - Quota limit
- Uses API keys to identify API clients

# <sup>REST</sup> Validation and Transformation

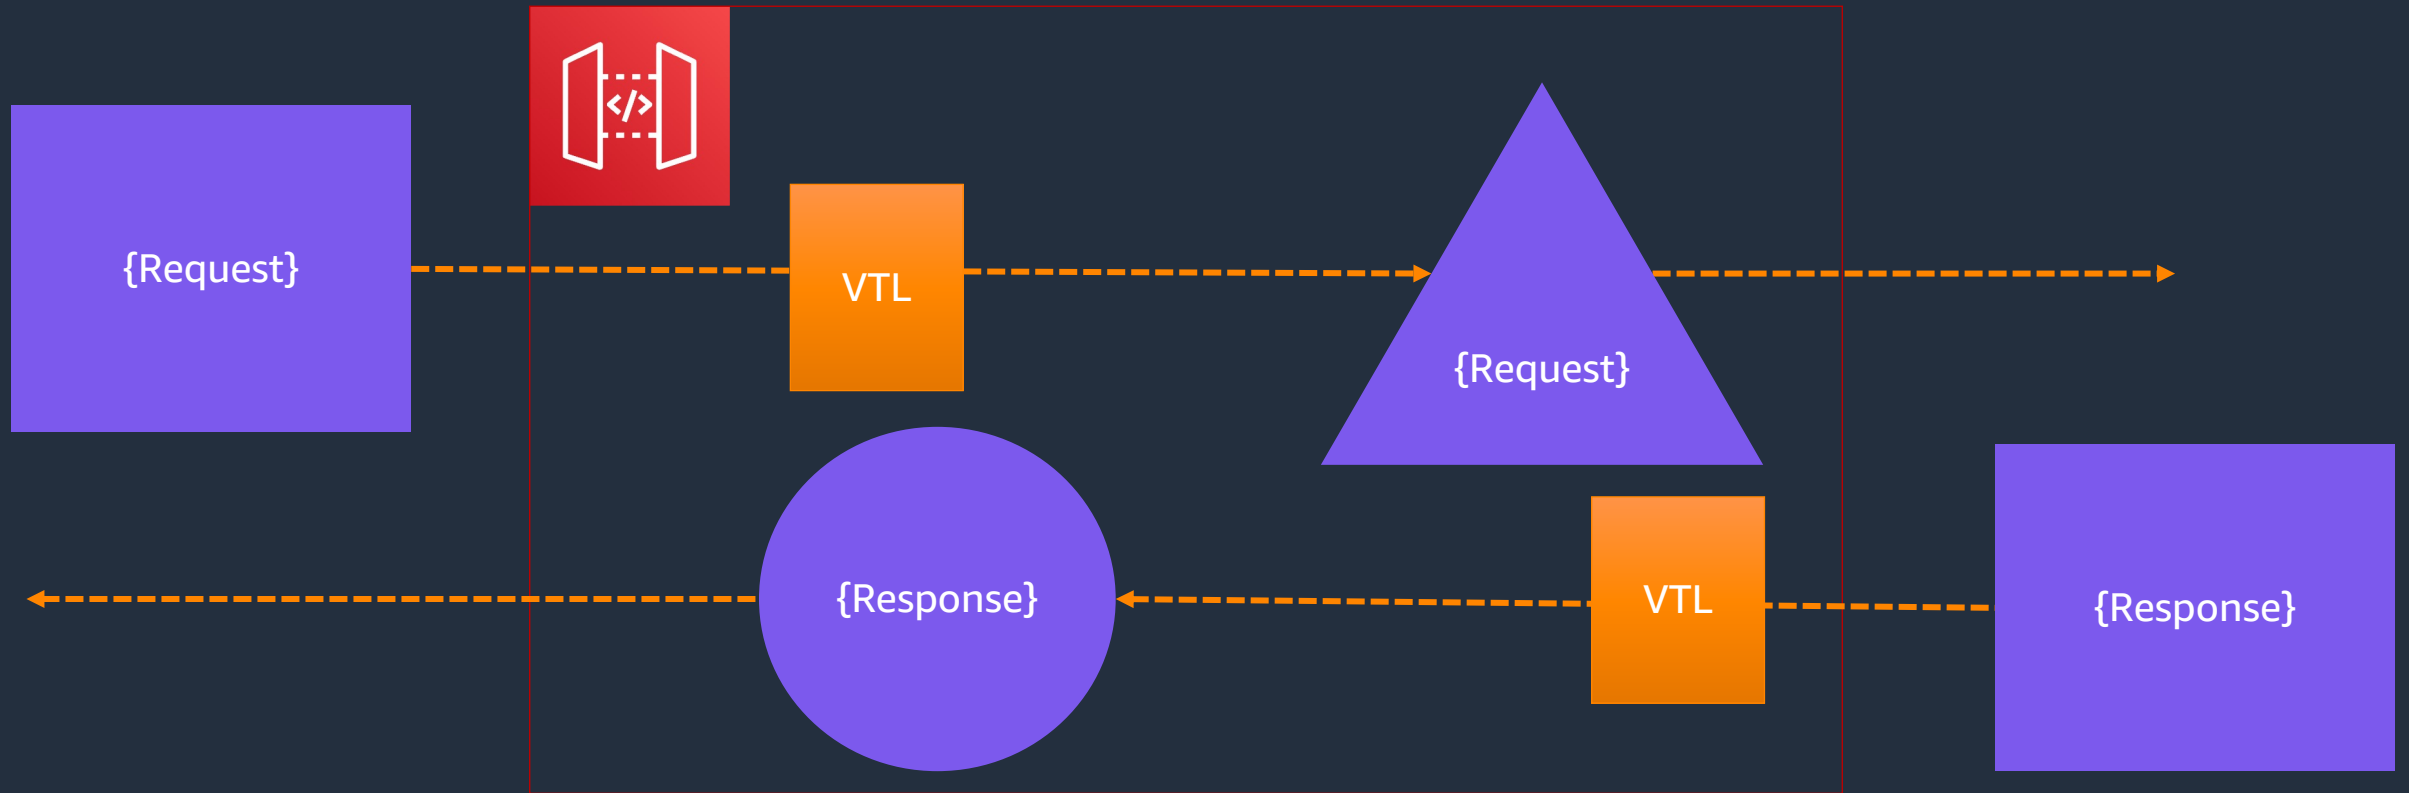




# REST Validation

- The required request parameters in the URI, query string, and headers of an incoming request are included and non-blank.
- The applicable request payload adheres to the configured JSON schema request model of the method.

# REST Transformation - Mapping



# Stages & Versions



# Stages

API Gateway enables you to set stage variables, allowing the same API to point to different backends.

Your APIs are versioned and can be rolled back.

- APIs are deployed to staging environments.  
You choose what to name them.
- For example, these environments:  
Dev (e.g., example.com/dev)  
Beta (e.g., example.com/beta)  
Prod (e.g., example.com/prod)

# Stages API Gateway



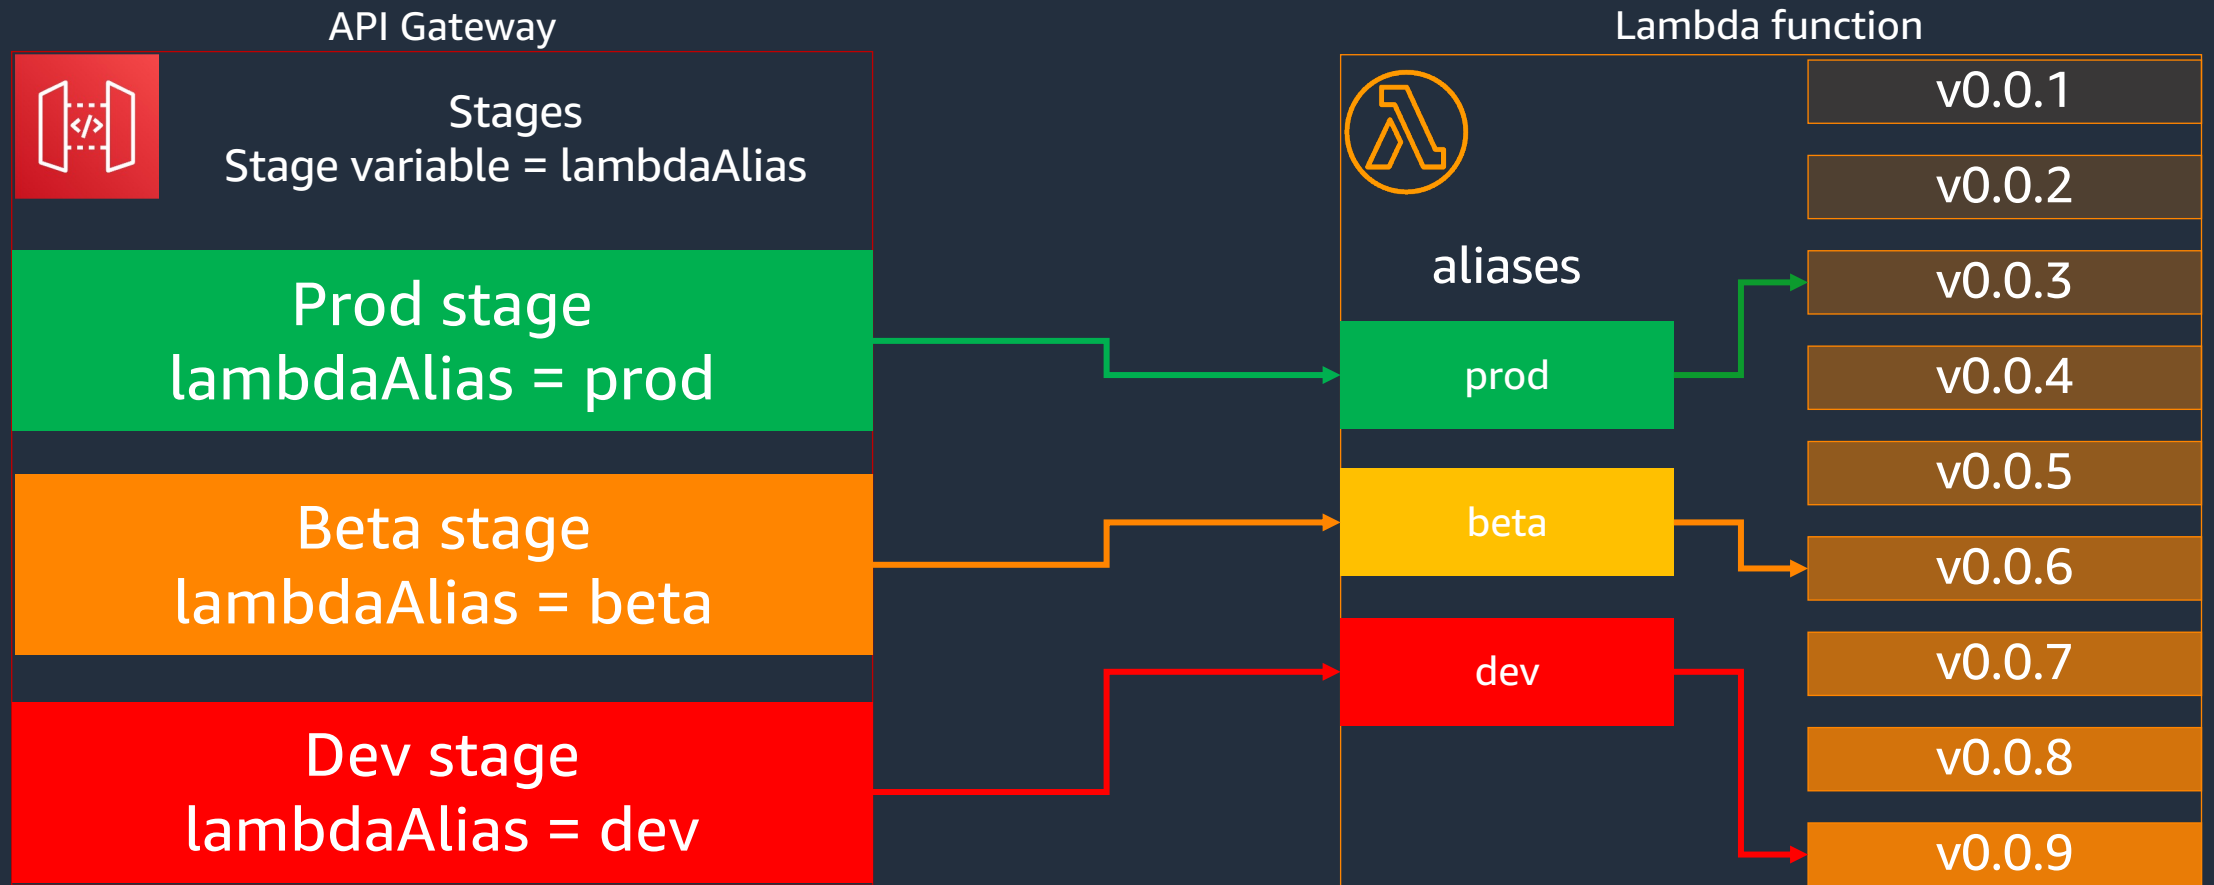
Stages  
Stage variable = lambdaAlias

Prod stage  
lambdaAlias = prod


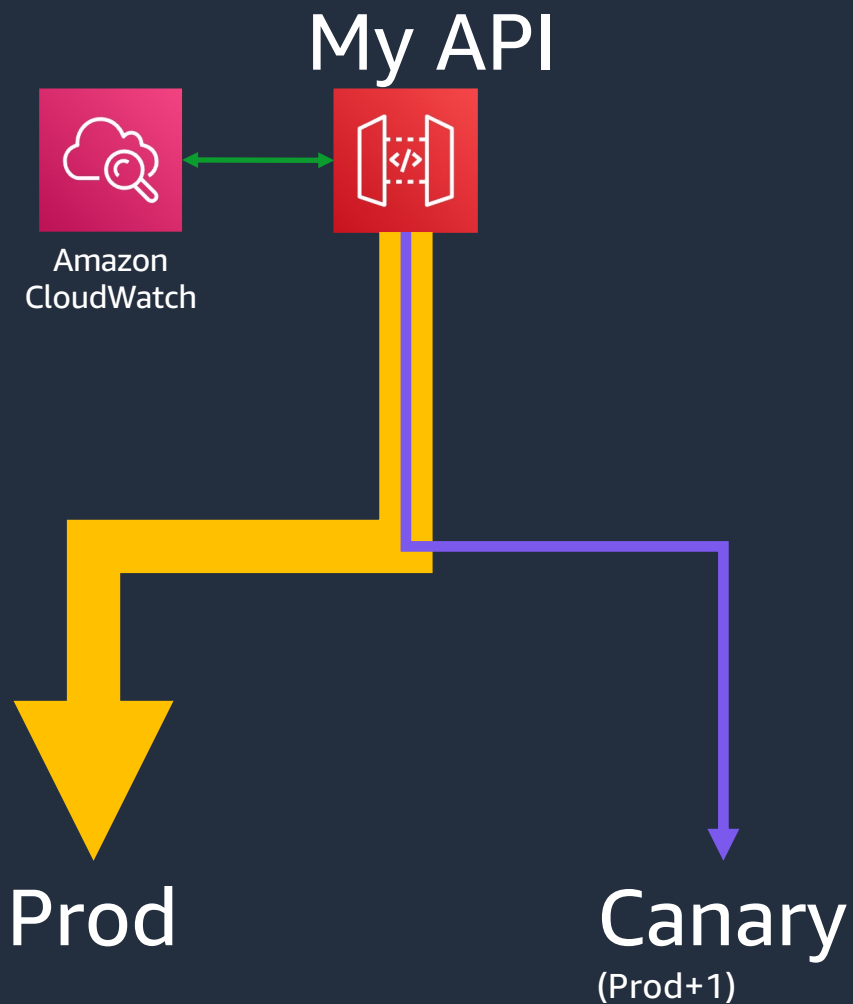
Beta stage  
lambdaAlias = beta

Dev stage  
lambdaAlias = dev

# Stages



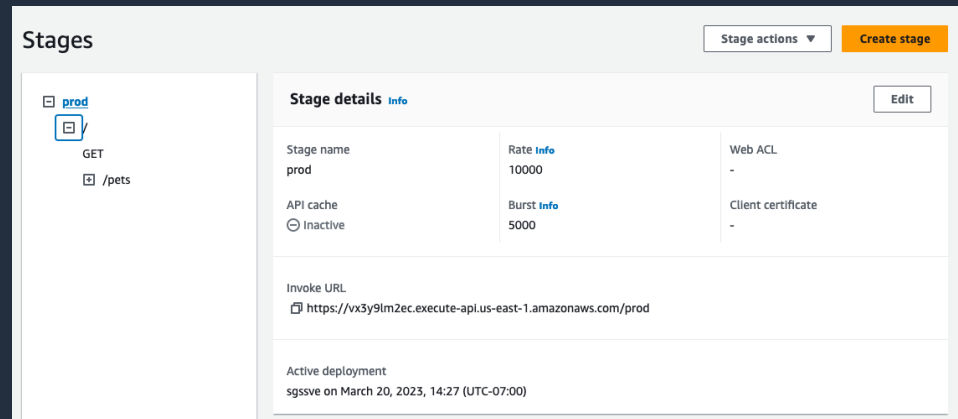
# REST Canary Releases



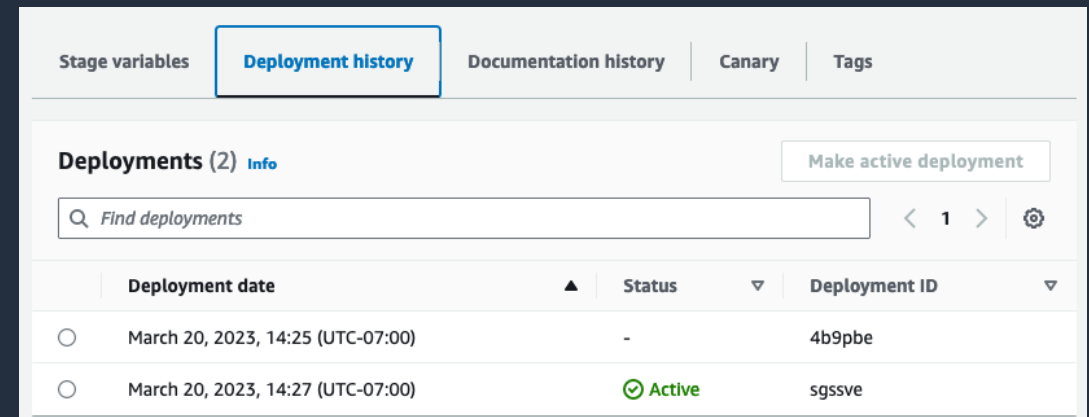
# REST API Versioning

We support versioning inside the API Console allowing to easily roll back to a previous API version and deploy it.

- Easily roll back to different snapshots of the same API (“Deployments”).
- Each Stage points to a Deployment. You can update the Stage to point to a previous Deployment.
- Permits Canary Deployments.



The screenshot shows the 'Stages' section of the AWS API Gateway console. On the left, a tree view shows the 'prod' stage under the 'GET /pets' endpoint. The main panel displays 'Stage details' for the 'prod' stage. It includes fields for 'Stage name' (prod), 'Rate' (10000), 'Web ACL' (-), 'API cache' (Inactive), 'Burst' (5000), 'Client certificate' (-), 'Invoke URL' (https://vx3y9lm2ec.execute-api.us-east-1.amazonaws.com/prod), and 'Active deployment' (sgssve on March 20, 2023, 14:27 (UTC-07:00)).



The screenshot shows the 'Deployments' section of the AWS API Gateway console. It features a tabbed interface with 'Deployment history' selected. Below the tabs is a search bar labeled 'Find deployments' and a 'Make active deployment' button. A table lists two deployments:

	Deployment date	Status	Deployment ID
<input type="radio"/>	March 20, 2023, 14:25 (UTC-07:00)	-	4b9pbe
<input type="radio"/>	March 20, 2023, 14:27 (UTC-07:00)	Active	sgssve



# Custom Domain Names



# Custom Domains

`https://12345.execute-api.us-east-1.amazonaws.com/prod/catalog`

The diagram illustrates the components of the AWS API endpoint URL `https://12345.execute-api.us-east-1.amazonaws.com/prod/catalog`. It uses green brackets to group parts of the URL and labels them below:

- API ID**: Points to the `12345` segment.
- Region**: Points to the `us-east-1` segment.
- Stage**: Points to the `prod` segment.
- Resource**: Points to the `catalog` segment.

# Custom Domains

~~https://12345.execute-api.us-east-1.amazonaws.com/prod/catalog~~

Base path mapping	API name	Stage
api-one	API1	prod
api-two/v2	API2	dev

***https://mydomain.com/api-one/products***

***https://mydomain.com/api-two/v2/catalog***



- Supports HTTP, REST, and WebSocket APIs
- SSL Certs managed through ACM
- Supports multiple APIs through multi level base path mapping

# Observability



# Metrics

## Built-in

### REST

API Calls Count, Latency, 4XXs, 5XXs,  
Integration Latency, Cache Hit Count,  
Cache Miss Count

### HTTP

API Calls Count, Latency, 4XXs, 5XXs,  
Integration Latency, DataProcessed

### WebSocket

Connect Count, Message Count,  
Integration Error, Client Error, Execution  
Error, Integration Latency

## Custom

Create Custom Metrics via Metric Filter  
out of logs

# Logging

## Execution Logs

Two levels of logging, ERROR and INFO

Optionally log method request/body content

Set globally in stage, or override per method

REST  
WS

Logs and tracing settings

By default, methods inherit the settings applied at the stage level. You can override the settings at the method level.

CloudWatch logs

Full request and response logs

☒ Detailed metrics

Each method will generate these metrics: API calls, Latency, Integration latency, 400 errors, and 500 errors.

## Access Logs

Customizable format for machine parsable logs

CloudWatch Logs OR Kinesis Firehose

REST

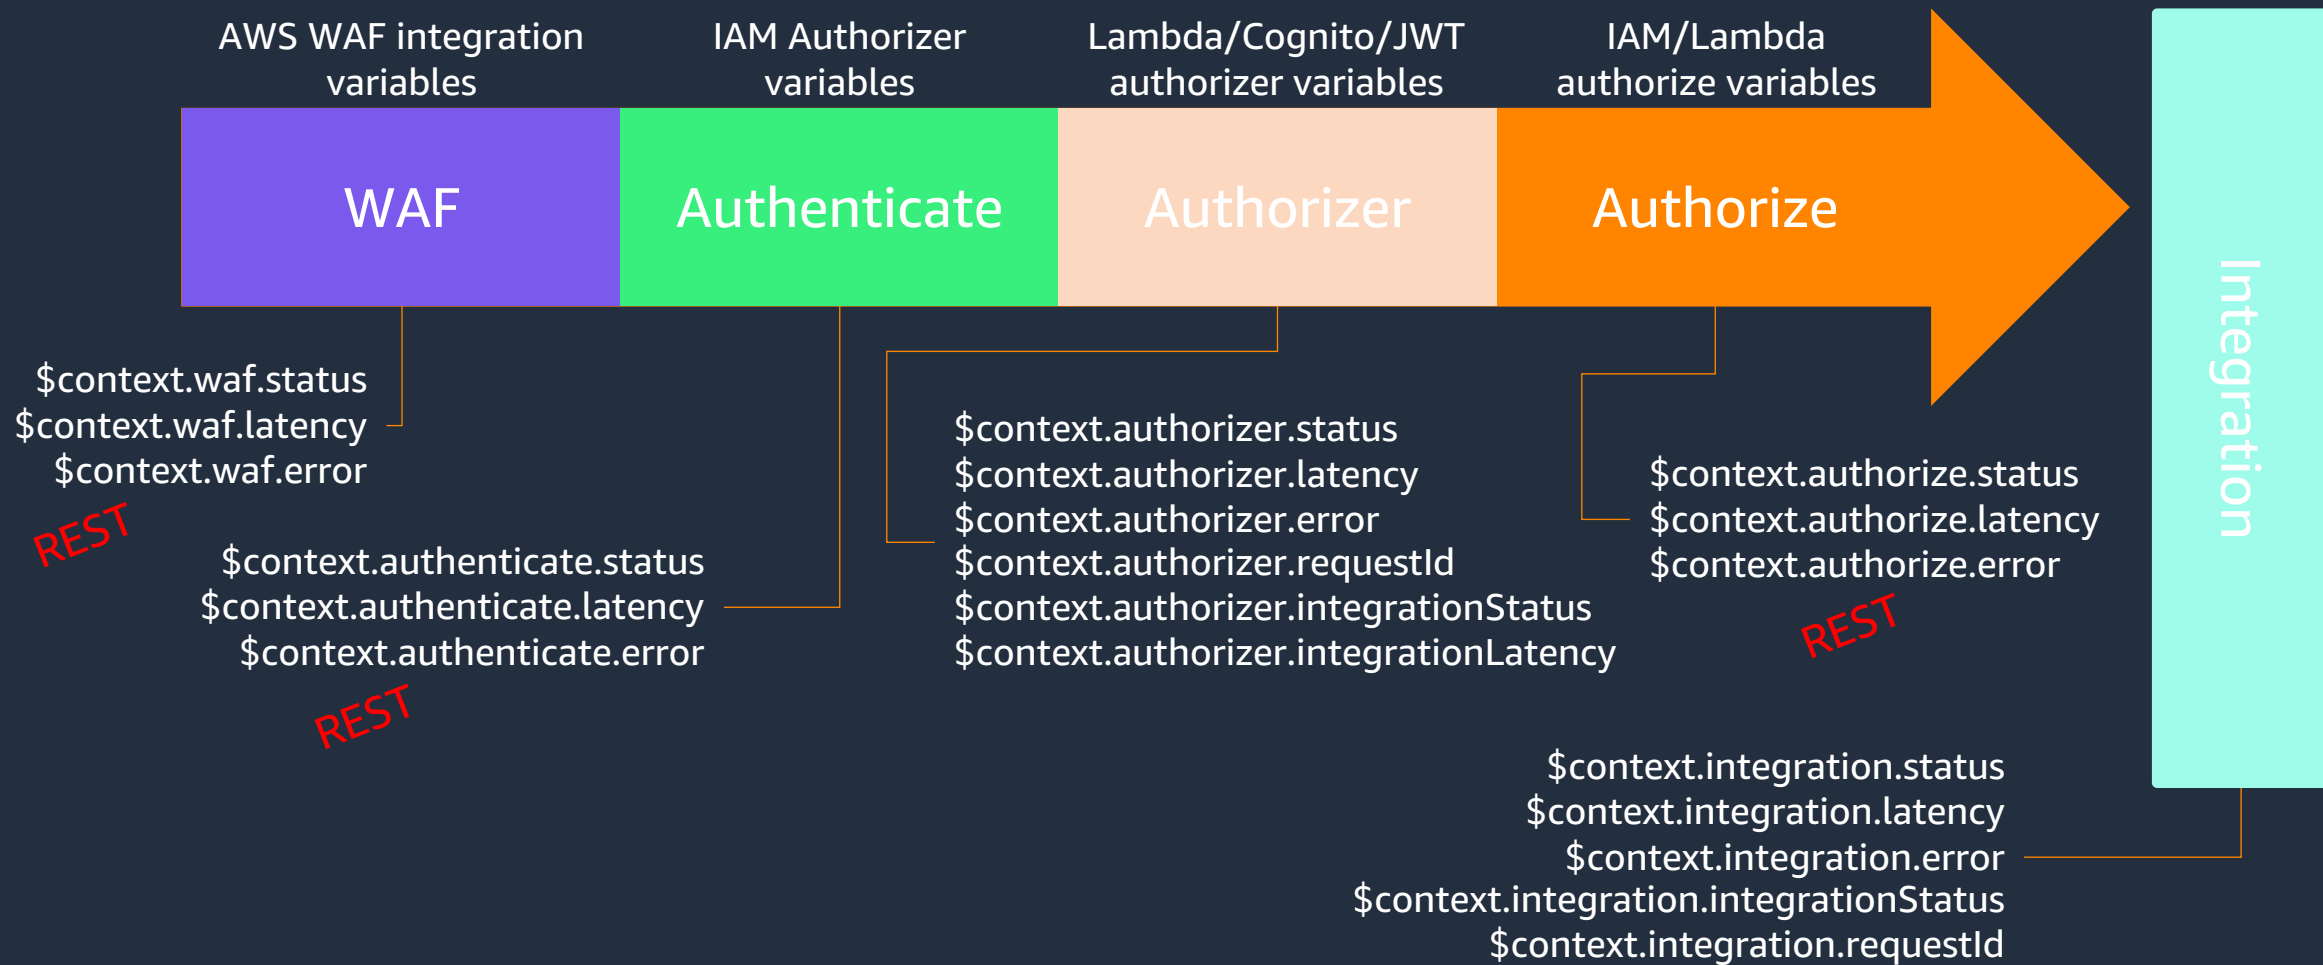
```
5c9-11e7-8228-318bf0a162b7) Verifying Usage Plan for request: 59b1
5c9-11e7-8228-318bf0a162b7) API Key authorized because method
5c9-11e7-8228-318bf0a162b7) Usage Plan check succeeded for AP
5c9-11e7-8228-318bf0a162b7) Starting execution for request: 59b1
5c9-11e7-8228-318bf0a162b7) HTTP Method: GET, Resource Path:
5c9-11e7-8228-318bf0a162b7) Method request path: {}
5c9-11e7-8228-318bf0a162b7) Method request query string: {}
5c9-11e7-8228-318bf0a162b7) Method request headers: {Accept=te
5c9-11e7-8228-318bf0a162b7) Method request body before transfo
5c9-11e7-8228-318bf0a162b7) Endpoint request URI: https://lambd
5c9-11e7-8228-318bf0a162b7) Endpoint request headers: {x-amzn-
5c9-11e7-8228-318bf0a162b7) Endpoint request body after transfor
5c9-11e7-8228-318bf0a162b7) Sending request to https://lambda.u
5c9-11e7-8228-318bf0a162b7) Received response. Integration later
5c9-11e7-8228-318bf0a162b7) Endpoint response body before tran
5c9-11e7-8228-318bf0a162b7) Endpoint response headers: {x-amzn
5c9-11e7-8228-318bf0a162b7) Method response body after transfo
5c9-11e7-8228-318bf0a162b7) Method response headers: {X-Amzn
5c9-11e7-8228-318bf0a162b7) Successfully completed execution
5c9-11e7-8228-318bf0a162b7) Method completed with status: 200
```

# Enhanced observability variables




status  
\$context.<phase>.latency  
error

# Enhanced observability variables





# Troubleshooting latency



Example of an IAM Authorized Endpoint with an AWS WAF in front of it

REST

# Tracing – X-Ray

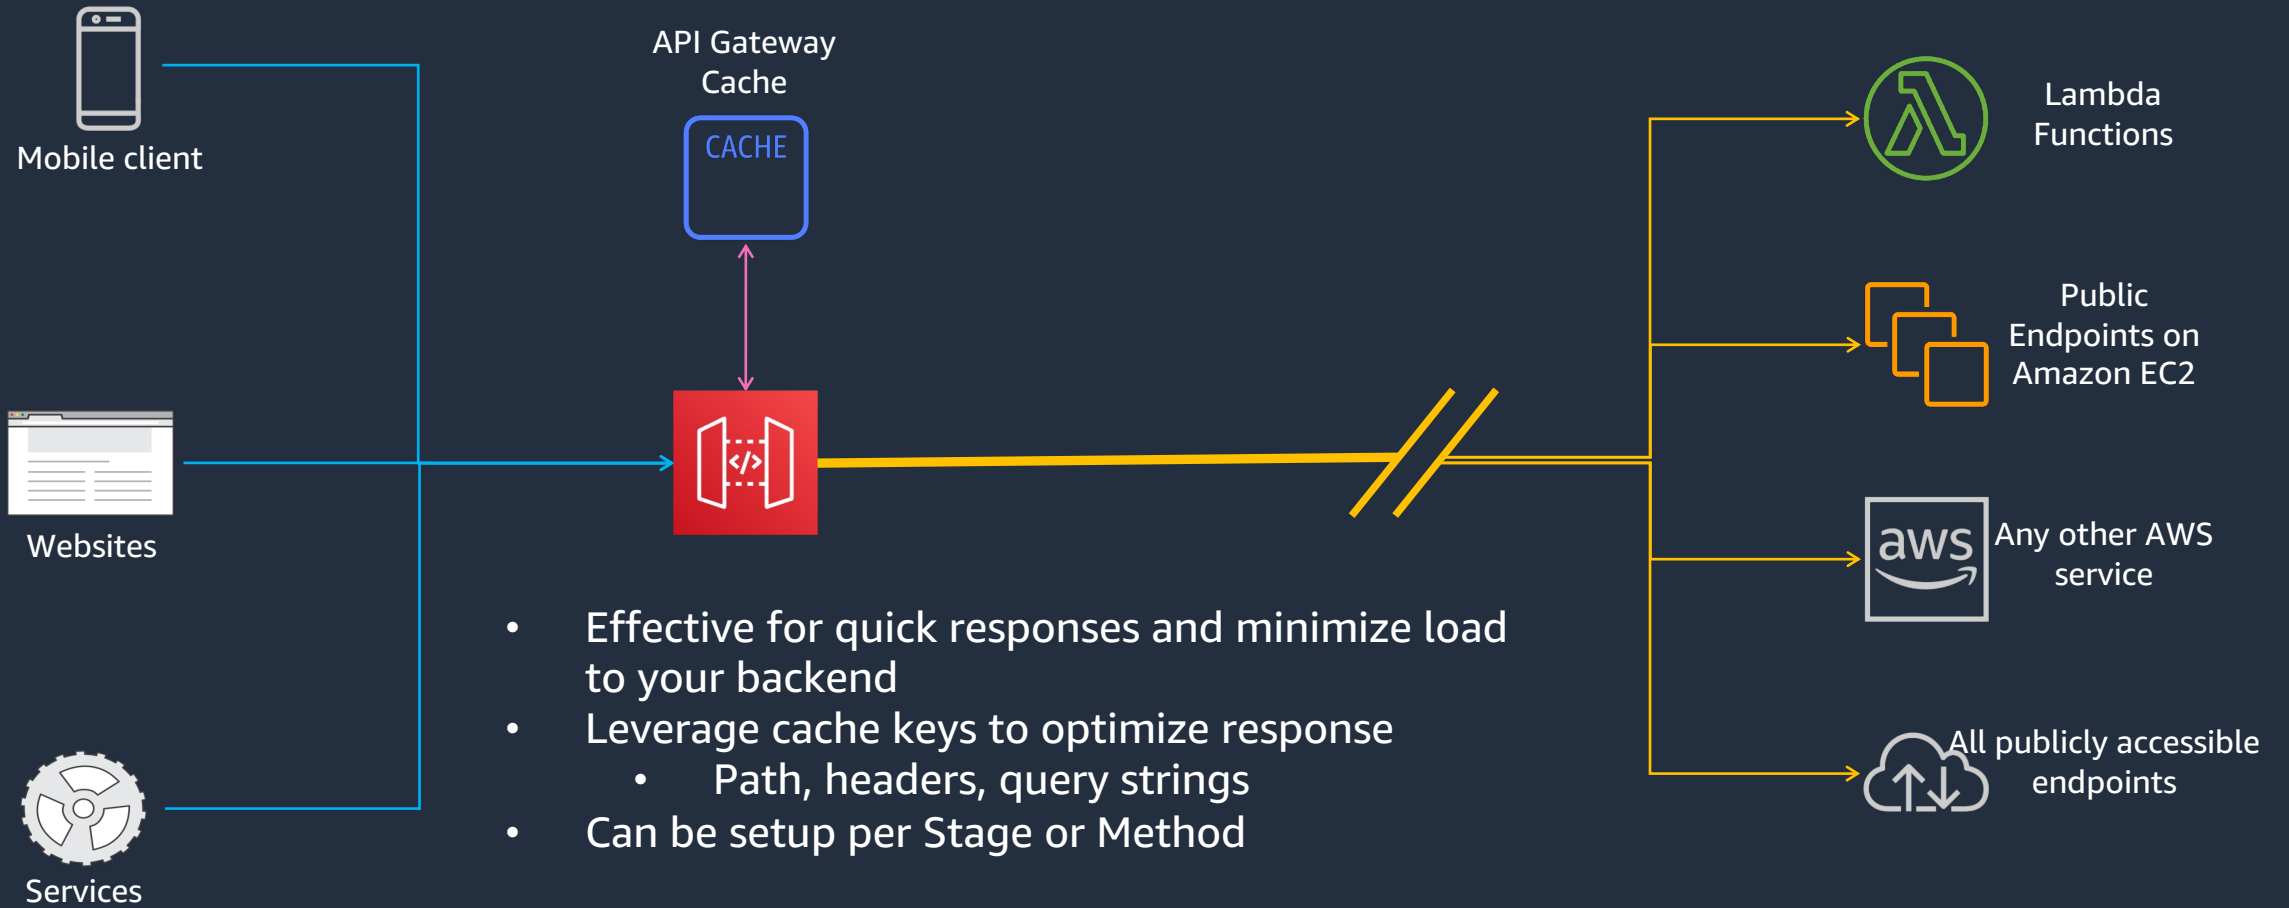
The screenshot displays the AWS CloudWatch Management Console's Service Map interface. The left sidebar contains navigation links for CloudWatch, Dashboards, Alarms (with a status summary showing 0 ALARM, 1 INSUFFICIENT, and 0 OK), Billing, Logs, Metrics, Events, Rules, Event Buses, ServiceLens, Service Map (highlighted), Traces, Container Insights (marked BETA), Resources, and Performance Monitoring. The main content area is titled 'CloudWatch > Service Map'. It features a time range selector (5m, 15m, 30m, 1h, 3h, 6h, custom) and buttons for 'Full screen', 'Map view', and 'List view'. Below these are search filters: 'Filter by X-Ray group' and 'Select a node', along with a 'View connections' button and a 'Requests mode' dropdown. The central part of the screen shows a service graph with nodes representing different AWS services and resources, connected by lines indicating the flow of requests. A 'Map legend' on the right explains the symbols: Fault (5xx) in red, Error (4xx) in orange, Throttle in purple, and Current alarms with a warning icon. It also defines the shapes for AWS services (rounded rectangle), AWS resources (oval), and unknown nodes (dashed oval), and notes that the size of the circles represents the number of requests. At the bottom of the graph area, a message states 'No node selected'. On the right side of the graph area, there are buttons for 'View logs', 'View traces', and 'View dashboard'. The footer of the console shows the URL 'https://eu-west-1.console.aws.amazon.com/cloudwatch/home?region=eu-west-1#servicelens:map' and copyright information for Amazon Web Services, Inc. or its affiliates.




# Other Features



# Caching



# OpenAPI for migrating APIs



# Pricing



# Pricing (eu-west-1)

- HTTP API requests: \$1.11 / Million (1M)\*
- REST API requests : \$3.5 / Million (1M) \*
- REST API cache: \$0.2/hour for 0.5GB ➡ \$3.80/hour for 237 GB \*
- WebSocket requests: \$1.14 / Million messages \* (1M)
- WebSocket connection minutes: \$0.285 / Million minutes (750,000)

\* Tiered pricing

*Free tier per month for 1 year*



# Important quotas





# Important quotas

- **Throughput:** 10,000 Requests/second
- **Max integration timeout:** *30 seconds*
- **Payload size:** *10 MB*
  
- **WebSocket connections:** 500 connections/sec
- **WebSocket connection duration:** *2 hours*
- **WebSocket message size:** *128 KB*

More: <https://docs.aws.amazon.com/apigateway/latest/developerguide/limits.html>



# Best practices

- Use HTTP APIs if existing features are sufficient
- Create an API per team/microservice and unify them with Custom Domain Name
- The console is for experimenting, use Infrastructure as Code:
  - SAM
  - CloudFormation
  - Terraform
  - Etc.

# Useful resources

## Security Overview of AWS API Gateway:

<https://d1.awsstatic.com/whitepapers/api-gateway-security.pdf>

## Serverless Lens:

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/welcome.html>

## Choosing between HTTP and REST APIs

<https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-vs-rest.html>

## Best Practices for Designing Amazon API Gateway Private APIs and Private Integration

<https://docs.aws.amazon.com/whitepapers/latest/best-practices-api-gateway-private-apis-integration/best-practices-api-gateway-private-apis-integration.html>



# Serverlessland.com



Home

Blogs

Videos

Learn

Events

About

Search here...



## Welcome to Serverless Land

This site brings together all the latest blogs, videos, and training for AWS Serverless. Learn to use and build apps that scale automatically on low-cost, fully-managed serverless architecture.

Learn More

